

Master Thesis

Nicola Covallero

May 16, 2016

Contents

Contents	1
1 Introduction	2
1.1 Problem Approach	2
1.2 Set Up	3
2 State of the Art Manipulation Planning	6
2.1 Segmentation	8
2.2 Review Pushing	8
2.3 Review Grasping	8
2.4 Usefull things	9
2.5 Usefull Concepts	9
2.6 Problem Description	9
2.7 Real Robot	9
3 Task Planner	10
3.1 State of the Art	10
3.2 Planner	11
3.3 Symbolic Predicates	12
4 Algorithm	16
4.1 Object Localization	16
4.1.1 Tabletop Object Detection	16
4.1.2 Object Segmentation	17
4.2 Predicates Computations	18
4.2.1 Tools	19
4.2.2 Predicate: <code>block_grasp</code>	19
4.2.3 Predicate: <code>on</code>	19
4.2.4 Predicate: <code>block_dir_i</code>	19
5 Implementation	20
6 Experiments	21
7 Further Work	24
Bibliography	25

1. Introduction

Robotic manipulation of objects is an increasing field of research which has struggled researchers from many years ago. In several industries environments robots can be easily programmed considering some standard objects, i.e. the manipulation is always the same and known, and usually robot operations avoid to deal in cluttered scene, for instance through a conveyor belt the objects are singulated from the others ones. In order to move the robots outside industries and bring them to home of people they need to be equipped with a particular intelligence and skills that are still at a first stage. In this thesis we are going to investigate a simple approach that tries to replicate the human reasoning step during the manipulation of objects of table clearing task. The main objective is trying also to manipulate the objects in a human style, that is paying attention that during the manipulation the collision between objects is avoid as much as possible. This is a behaviour that all we do, of course with fragile objects, but also with objects that have no problem regarding collisions.

Next, in this chapter a brief introduction of the problem and the approach used is commented, and the experimental set up as well since it will be useful to understand some techniques and choice for the algorithm.

1.1. Problem Approach

In this section the approach to resolve the planning problem is described. All the strategy is inspired to a human like solution. A human to resolve such a task would use three main actions: grasp an object with one hand or both, when it is not possible to grasp an object because other object impeded the human to put the hands in the proper way to grasp a desired object he/she has to interact with the objects in order to achieve a configuration of the objects that can be suitable to grasp an object, and this is achieved through pushing or dragging actions. Dragging an object is a very complex action which will need the robot to lay the end effector on the top of an object and make enough pressure on the contact in order that the friction between the table and the object is minor than the friction between the object and the robot's end effector. This is a very hard action which would imply an implementation of a reliable controller, because the goal is not only to move the object but also not to ruin it.

For these considerations the main actions the robot has to use in order to interact with the objects are:

- Grasping
- Pushing

Grasping is the most important action since it lets to take an object from the pile of objects and drop it somewhere, for example in a bin, clearing in this way the table.

There exist different works facing the same task by focusing only in grasping. The pushing becomes useful considering the problem that two adjacent objects could not be grasped if they are so close such that the robot's gripper, when attempting to grasp an object is going to collide with an adjacent object, making such an object ungraspable. From this consideration is necessary the pushing action, in order to separate adjacent objects which could mutually exclude themselves to be grasped. Moreover the robot's intelligence is enhanced by a planning system which will return a sequence of actions in order to achieve the goal to clear the table.

1.2. Set Up

In order to make the understanding of the approach used the name the set up of the environment the robot will work in is presented. The robot used is a Barret WAM arm, which is a 7 degree of freedom (DoF) manipulator arm shown in Figure 1.1b.

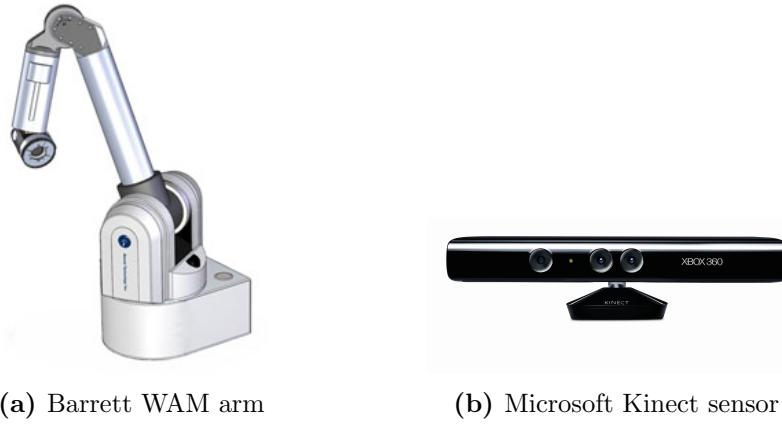


Figure 1.1: Robot and vision sensor.

This robot is characterized by a low inertia of the end effector thanks to the kind of its actuators. The WAM is noteworthy because it does not use any gears for manipulating the joints, but rather a cable drive, so there are no backlash problems and it is both fast and stiff. Cable drives permit low friction and ripple-free torque transmission from the actuator to the joints.

To manipulate the objects the manipulation skills of the robot are enhanced by a gripper designed in the IRI institute driven by Dynamixel motors. Such a gripper is depicted in Figure 1.2 from several point of views.

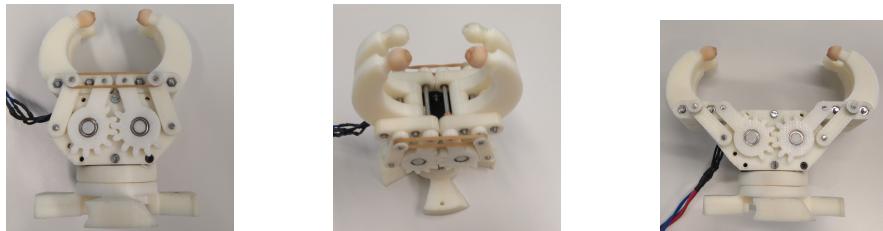


Figure 1.2: Gripper used for the experiments

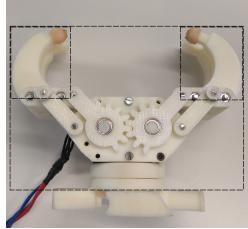
For the task planner, as the reader will see in chapters 3 and 4, the model of the gripper will be an important resource in order to find a feasible sequence of action that can free the table.

The gripper will be modelled measuring some principal elements such as: finger's width and height, gripper's width, height and deep, and the distance between the

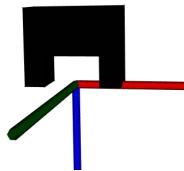
fingers when the gripper is opened or closed. The modelling procedure is depicted in Figure 1.4. The resulting model is a simple polygonal mesh which includes all the important geometric information of the gripper.

Such a simple model allows the collision algorithm commented in chapter 4 to check for collision in just few millisecond. A more detailed and complex model would only lead as benefits a higher precision, precision that in this case is not needed, and it will slow down algorithm because the collision checking procedure would be more computationally expensive. The gripper will be mounted in the end effector of the robot as shown in Figure 1.3.

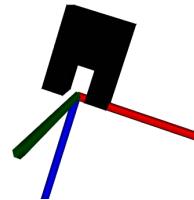
Figure 1.3: Gripper and WAM.



Elements measured



Opened gripper mesh model

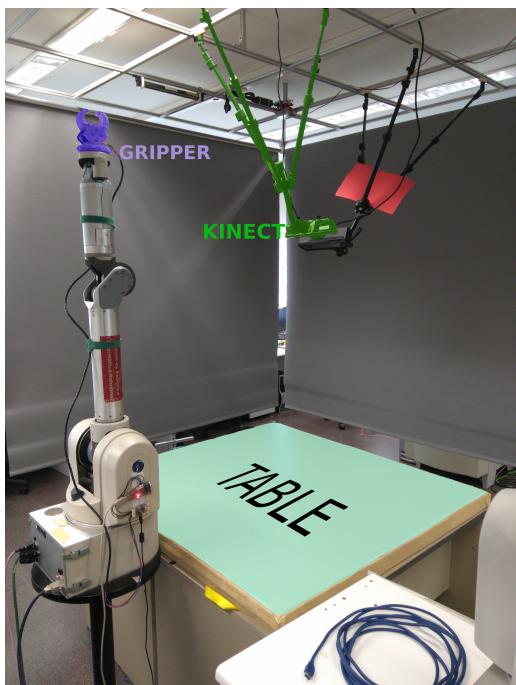


Closed gripper mesh model

Figure 1.4: Modelling of the gripper: the fingers width, fingers height, the height of the whole gripper, its width and its deep have been measured. The red, green and blue axis are respectively the x, y and z axis.

The scenario the robot is going to work in is presented by a table, and the objects will lay on top of it. In Figure 1.5a the main elements of are highlighted. The WAM arm is in a fixed position with respect the table and the Kinect camera is located on top of the table pointing down. The Kinect is calibrated and the fixed transformation between the Kinect's frames and the base frame of the robot is known, so all the points measured by the Kinect can be expressed in coordinates with respect the robot's frame. An example of a cluttered scene the robot is going to deal with is depicted in Figure 1.5b, and the same scene seen by the Kinect is shown in Figure 1.5c. With respect the Kinect's view the robot is located at the top of the image, outside the point of view of the Kinect in order to avoid to detect the arm as an object. If the robot arm would be present in a depth image we should apply some algorithms to identify the robot arm, which could also produce important occlusions, that is hide some objects from the Kinect's view. In order to avoid all this, the depth image considered is the last one provided by the Kinect when the robot is no more in the Kinect's view (these poses are known a priori).

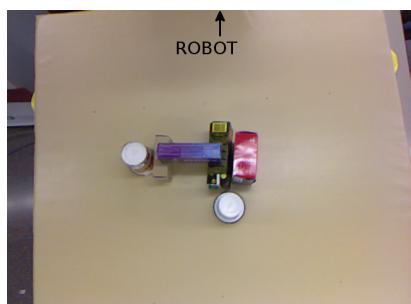
Concerning the example of Figure 1.5b it can be appreciate that the sequence of operation needed in order to clean all the table with a careful manipulation is quite complex, and due to the gripper geometry, and to the grasping poses we choose (section ??), the robot will need to grasp first the purple top object, or the white cup, then move one box in other to grasp the one of its side. All this sequence of actions will be decided by the robot.



(a) Principal elements of the experimental set up.



(b) Example of a cluttered scene the robot is going to work with.



(c) Kinect's view.

Figure 1.5: Experimental set up, and example of a cluttered scene the robot is going to interact with.

2. State of the Art Manipulation Planning

cite David's works

still to add reference to the ones in the pushing.odt file

Many manipulation planning approaches (see [27] for an overview) assume that the task can be treated as a geometric problem, with the goal to place the objects in their desired positions. Planning is essentially done with a mixture of symbolic and geometric states, this requires a set of symbolic predicates that correspond to geometric relationships. In particular to the best of my knowledge, all the planners for cluttered scene are based on a mixture of symbolic and geometric predicates, or based only on geometric ones.

Dogar and Srinivasa [18] proposed a framework for planning in clutter scenes using a library of actions inspired by human strategies. They designed a new planner that decides which objects to move, the order to move them, where to move them, and it chooses the manipulation actions to use on these objects, and accounts for the uncertainty in the environment all through this process. The planner first attempts to grasp the goal object, than if it is not possible it identifies what is the object that obstacles the action, then such an object is added to a list of the objects that have to be moved. They moved an object in whatever position such that makes the goal feasible. Despite this, such a planning strategy can be used only in the design of a new planner, which is something we want to avoid.

To grasp they used the Push-grasping action [17], which is a robust way of grasping objects under uncertainty. It is a straight motion of the hand parallel to the pushing surface along a certain direction, followed by closing the fingers. For the pushing directions they used a resolution of $\pi/18$ rad (i.e. 36 different directions) and use a predefined set of 9 hand aperture values. However their cluttered scenes is intended to be a scene with separated and well known objects.

In [21] the authors proposed to apply a linear temporal logic (LTL) planner to a manipulation planning task in cluttered scenes, but it suffers from poor runtime and the LTL specification is complex. Moreover to the scope of this work the temporal specification is not needed since it is indirectly encoded in the clutter scene composition.

An interesting planner which mixes symbolic and geometric predicates is aSyMov [12]. It is probably the most well known planner which a task to move objects considering how their location will then affect the possible robot's paths. To do that they combine a symbolic planner with a probabilistic roadmap [39] for geometric planning.

A recent alternative proposed by Mösenlechner and Beetz [33] is to specify goals symbolically but evaluate the plan geometrically. The idea is to use a high-fidelity

physics simulation to predict the effects of actions and a hand-built mapping from geometric to symbolic states. Planning is conducted by a forward search, and the effects of actions determined by simulating them, and then using the mapping to update the symbolic state. This could be look promising to the scope of this thesis but still implies the design of a new planner.

In [29], Lonzano-Peréz et al. They use a simple task-level planner, in which operators are described with two types of preconditions: symbolic and geometric. They proposed a strategy for integrated task and motion planning based on performing a symbolic search for a sequence of high-level operations, such as pick, move and place, while postponing geometric decisions, based on the CSP (Constraint Satisfaction Problem) technique. **Their technique allow working directly with state of the art planners.**

In [13] the authors address a problem similar to the one of this thesis. Pushing presents itself as a more intricate problem than grasping, a for grasping, after the object has been grabbed, it becomes a matter of solving a geometric problem. However, pushing an object carries uncertainty to the resulting state of the object. It is, nonetheless, important to be able to incorporate these two actions in the planning. For pushing they used a similar strategy to [32]. The planning algorithm also incorporates grasping, which, together with push actions, can perform a vaster array of tasks. It uses the concept of **reachability** [41] for each kind of action to quickly exclude impossible poses of the gripper of the planning stage, creating a reliable plan suitable for real-time operation. The authors model the planning problem through MDP (Markov Decision Process), discretizing the world in grid cells and assigning to each one a push and grasp vector defined by the reachability concept. Moreover each object is also classified through a simple primitive shape in order to have a proper map of pushing and grasping for each kind of considered object.

The majority of the state of the art for task planning in cluttered scene is focused on designing a new planner, why in this thesis state of the art planners ready to use want to be used.

Learning in Task Planning Machine learning techniques also have been applied to task planning. To plan complex manipulation tasks, robots need to reason on a high level. Symbolic planning, however, requires knowledge about the preconditions and effects of the individual actions. In [9] the authors proposed a practical approach to learn manipulation skills, including preconditions and effects, based on teacher demonstrations. It is difficult to solve most real world manipulation tasks by reasoning purely in terms of low-level motions due to the high-dimensionality of the problems. Instead, robots should reason on a symbolic level and appropriately chain the learned actions to solve new tasks. Such a planning step, however, requires knowledge of the important preconditions and effects of the actions. With few demonstration the authors proposed a method to teach the robot the precondition and effects of individual actions. Their method furthermore enables the robot to combine the learned actions by means of planning to solve new tasks that are more complicated than the learned individual actions. They focused their work on a class of manipulation actions where the preconditions and goals depend on spatial or geometrical constraints between the involved objects. They describe a scene as a collection of features and during the action demonstration the algorithm looks for the features that have a small change to identify the ones that can be predicates or effects of the considered action. They use an heuristic metric to define the variation of the features. The problem of this approach is that it moves the problem of

identifying the right predicates to describe the problem in identifying right features that let a correct learning of the predicates. **We could consider this as further work.**

In [16] [15] the authors propose an approach for planning robotic manipulation tasks which uses a learned mapping between geometric states and logical predicates. The planner first plans symbolically, then applies the mapping to generate geometric positions that are then sent to a path planner. They try to fit a probability distribution function to the geometric states to map them to a symbolic state. This work sounds promising but the cluttered scene this thesis is going to treat is quite complex to be treated with such a method.

Interpreting the scene Focusing on symbolic planning, the research group of Artificial Intelligence and Robotics Laboratory of Istanbul Technical University, published some interesting researches suitable to the aim of this thesis. In [19] [34] [20] the authors propose a system which combines 3D recognition and segmentation results to create and maintain a consistent world model involving attributes of the objects and spatial relations among them. Unknown objects are modelled by using the segmentation output to determine their sizes and considering similarities with existing models to determine their shapes and colors. Then, these models are also stored as templates to be used for recognition along with the extracted attributes. They focused on the extraction of size, shape and color attributes as well as the following unary and binary spatial relations: *on*, *on ground/on table*, *clear* and *near* for object manipulation scenarios. These predicates are generated and updated over time to maintain a consistent world model. **For the aim of this thesis this paper is of particular interest because let to compute nicely some useful symbolic predicates** The *on* relation for a pair of objects is determined by checking whether their projections on the table plane overlap. For each pair of objects the *near* relation is determined by comparing the distance between the centers of these objects in each dimension with the sum of the corresponding sizes in that dimension.

Check fro backtracking techniques in planning, way to evaluate a plan

2.1. Segmentation

[26] interact with pile of objects with the aim to complement the segmentation algorithm for table clearance tasks.

2.2. Review Pushing

Review of pushing techniques

- [23] object singulation
- [24] object interactive perception and the same author did a work similar [25]
- [30]

2.3. Review Grasping

Review of grasping techniques

In [42] they presented a uniform deterministic sequence and sets of sample over 2-sphere and SO(3).

In [14] paper of grasping
ROS object_manipulator

2.4. Usefull things

Temporal filtering to reduce the noise of the kinect [34].

2.5. Usefull Concepts

Closed world assumption [37] can be defined as having complete knowledge about the world, that is, the numbers and the attributes of all objects are known apriori.

2.6. Problem Description

2.7. Real Robot

3. Task Planner

In this chapter the general framework adopted is discussed, proposing a suitable task planner. After the review of the current state of the art of task planners, a proper planner is chosen and then a suitable description to the table clearing problem is discussed.

3.1. State of the Art

As already seen in chapter 2 there exist different kind of planners and they can be grouped in three main categories:

1. classical planners
2. hierarchical planners
3. probabilistic planners

Classical planners, as suggested by the name, are the more classical and easy to use. They are characterized by environments which are fully observable, deterministic, finite and static (changes happen only when the agent acts) and discrete (in time, actions, objects..) [38]. A deterministic problem is generally formulated as a 6-tuple $\Pi = \langle S^d, s_o^d, G^d, A^d, T^d, c^d \rangle$ [28], where:

- S^d is a finite set of states;
- $s_o^d \in S^d$ is an initial state;
- $G^d \in S^d$ is a goal state;
- $A^d(s)$ is a set of applicable actions for each $s \in S^d$;
- $T^d(a, s) \in S^d$ is a deterministic transition function for all actions $a \in A^d(s)$ and states $s \in S^d$;
- $c^d(a)$ is the cost to apply action a .

The solutions, or trajectories, τ_i are sequences of actions applicable from the initial state until the goal state. The cost of a trajectory $C(\tau_i)$ is the sum of the cost of the actions of the trajectory $\sum_{a \in \tau} c^d(a)$. The optimal solution is the solution with less cost: $\tau^* = \min_{\tau_i} C(\tau_i)$. A very well known classic planner is the Fast Downward planner [22].

Hierarchical planning, also called *Hierarchical Task Network*(HTN), works in a similar way to how it is believed that human planning works [31]. It is based on a reduction of the problem. The planner recursively decomposes tasks into subtasks, stopping when it reaches primitive tasks that can be performed directly by planning

operators. In order to tell the planner how to decompose nonprimitive tasks into subtasks, it needs to have a set of methods, where each method is a schema for decomposing a particular kind of task into a set of subtasks [1]. For this kind of planning technique a well known planner is JSHOP2 [4]. **Probabilistic planning** is a planning technique which consider that the environment is not deterministic but probabilistic. So the actions have a certain probability to obtain a certain state, and given an initial and final state the planner finds the solution path with the highest probability. A well known example on which this kind of planners are build on is the Markov Decision Process. A probabilistic problem is generally formulated as a 6-tuple $\Pi = \langle S, s_o, G, O, T, A \rangle$ [28], where:

- S is a finite set of states;
- $s_o \in S$ is the initial state;
- $G \in S$ is a goal state;
- O is the set of outcomes, the probability of $o \in O$ is $Pr(o)$;
- $T(o, s) \in S$ is a (total) deterministic transition function for all outcomes $o \in O$ and states $s \in S$;
- $A(s)$ is a set of applicable actions for each $s \in S$, coupled to a function $out(a) \subseteq O$ mapping each action to a set of outcomes in such a way that
 - each outcome $o \in O$ belongs exactly to one action $act(o)$;
 - $\sum_{o \in out(a)} Pr(o) = 1$ for all a .

In this case the optimal solution is the one with the highest probability. In this category two famous probabilistic planners are Gourmand [2] and PROST [3].

3.2. Planner

The problem this thesis is facing could be resolved by several approaches by using planners from all the categories. We have already seen that such a problem involves geometric constraints and those cannot be considered directly by the planner using a ready to use state of the art planner, that would imply a designing of a new planner. Since the aim of this work is not to design a new planner but to resolve the table clearing task through already existing planners the problem has to be cast in a way to be manipulated by existing planners. This easy way involves working with symbolic predicates, symbolic predicates are predicates which can be true or false, and they will be introduced more in detail in the next sections.

The problem moreover involves a big amount of uncertainty due to the interaction of the robot with the environment. When the robot will interact with the pile of objects it is very hard to predict correctly the position of the object in the next frame, that is the next state, this is a crucial problem which should be considered. With a probabilistic planner, the planner will take into account what object has been moved and it will update the state obtaining a set of states, each one with a certain probability, and the returned plan, or solution, is the one with the highest probability. The probability also has to be modelled and it is highly depending on the form of the object, which is also hard to predict. An other way to face the problem is to replan at each frame, that is after each interaction of the robot with the pile of objects, or whenever the current state deviates from the expected one,

generating a new trajectory from the current state to the goal. The plan's actions are considered deterministic, and the only useful action is actually the first one of the plan, after its execution the system replans again. Little et al. discussed in [28] the problem of when is more useful the probabilistic planning with respect a simple replanning system. They defined a the concept of *Probabilistic Interesting Problem* with the following definition:

A probabilistic planning problem is considered to be *probabilistically interesting* if and only if it has all of the following structural properties:

- there are multiple goal trajectories;
- there is at least one pair of distinct goal trajectories, τ and τ' , that share a common sequence of outcomes for the first $n - 1$ outcomes, and where τ_n and τ'_n are distinct outcomes of the same action; and
- there exist two distinct goal trajectories τ and τ' and outcomes $o \in \tau$ and $o' \in \tau'$ of two distinct actions $a = act(o)$ and $a' = act(o')$ such that executing a strictly decreases the maximum probability of reaching a state where a can be executed.

They assert that unless a probabilistic planning problem satisfies all of the structural conditions in this definition, then it is inevitable that a well-written replanner will outperform a well-written probabilistic planner. Moreover the authors do not negate the possibility that a deterministic replanner could perform optimally even for probabilistically interesting planning problems.

To conclude, a replanner would make more probabilistic problem practically solvable. In the other hand it suffers to be less robust than a probabilistic one.

Taking into account such considerations, in order to face simply such a difficult problem, the problem has been thought to be solved thanks a deterministic replanner although it is *probabilistic interesting*. The choice also was guided by the difficulty to model the probability distribution of the actions which depends on the particular shape of the objects the robot has to interact with.

Between a classic planner and an heuristic one there is not much difference for the aim of this work. The planner chosen to develop the work is the **Fast Downward** planner [22][5], a very well know classic one. The advantage of this planner is its wide documentation with respect other planners, which makes it easier to use and it also has a wide community.

3.3. Symbolic Predicates

One the planner has been chosen, a task planning problem has to be formulated accordingly to the capabilities of such a planner. The *Fast Downward* planner needs the problem to be formulated in the common format of PDDL (*Problem Domain Description Language*) [7].In this section the symbolic predicates that have been considered in order to resolve the problem are described.

The task this thesis is going to resolve is a common task done by humans who thinks in order to find a feasible sequence of actions. Such a sequence is normally composed of actions that avoid the collision between the manipulated object and the other ones, whenever possible. To do this we, as human, think on what is going to happen if we manipulate an object in a certain way, and what objects can be manipulated at the current instant and what objects need to be moved in order to interact with a certain one.

The design of the problem has been inspired on such a reasoning way. The reasoning part is entrusted to the planner, which needs a coherent problem description through symbolic predicates. How described in the introduction, the system will be able to perform two types of actions: **pushing** and **grasping**. Grasping action is a necessary action in order to grasp an object and drop it in a bin, while the pushing action is an auxiliary action which has the aim to move an object in order to collocate it in a pose that do not obstacle the solution of the plan. In fact it happens that an object cannot be grasped because there is an adjacent one that impedes such an action.

The symbolic predicates are designed accordingly to the available actions trying to answer the following questions:

- When cannot an object be grasped?
- When can a object be pushed? In which direction?

Grasping Action An object cannot be grasped always but only if the grasping action is not impeded by other objects, to do this the following things have to be checked:

- no objects have to stand on top of it
- there are no adjacent objects that impede the grasping action

We don't want to grasp an object which has another one on top of it for one main reason: during the grasping action such object will fall corrupting the scene and such behaviour is undesired, in fact a human will never grab an object at the bottom of another one without first grabbing the one on the top. To include such an information in the problem description the predicate **on** is created, in particular it is defined as (**on o1 o2**) meaning that object labelled as **o1** stands on top of **o2**.

Once we are sure the current object has not objects on top of it we have to check if it can be grasped, that is, given a grasping pose we have to check if that pose is feasible or if it is impeded by other objects. To include such an information in the problem description the predicate **block_grasp** is created. In particular it is defined as (**block_grasp o1 o2**) meaning that object labelled as **o1** impedes object **o2** to be grasped.

Pushing Action In certain situations when the grasping is not possible because an object impedes another one to be grasped one of both has to be moved. Firstly, being observant of the philosophy to avoid manipulation that can damage the objects, we considered that only the objects that stand on the table can be pushed, that is an object on top of another one will not be pushed. And, how discussed for the grasping action, also an object that has on top of itself another one will not be pushed. Second the possible directions along with an object can be moved have to be decided. Such directions are taken accordingly to the shape of the object, in particular to its principal directions. Moreover the pushing direction is considered to be parallel to the table plane, as a normal human would be. Therefore the principal direction is projected onto the table plane, and also the orthogonal direction to the principal one is considered, having in total 4 possible ways to push an object, 2 ways per direction. To the best of my knowledge there are no studies about how pushing an object on a proper way, or direction, in order to make it follow a desired trajectory.

Having the direction each object can be pushed along with, the next step is to understand if the object can be pushed along those directions. An object cannot be pushed in a certain direction if there is another one that blocks the current object. For each direction the object is therefore translated in the pushing direction and checked for collision with other objects. In this way it is possible to get what is the object which impedes the considered one to be moved along the considered direction.

To include such an information the predicates **block_dir1**, **block_dir2**, **block_dir3** and **block_dir4** are created. In particular, they are defined as $(\text{block_dir}_i \text{ o1 o2})$ meaning that the object labelled as **o1** is impeding object **o2** to be moved along direction dir_i .

Effects of the action The effect of the action has to change the current predicates of the problem in order to update coherently the state. For the *grasping* action whenever an object is grasped, since we are working in a deterministic scenario, the object is supposed to be grasped successfully and drop in the bin. If such an object was on top of other ones or it was impeding some objects to be pushed in some direction those predicates have to be updated coherently. Moreover to consider that the object is no more in the scene the predicate **grasped** is added. In particular (grasped o1) means that the object labelled as **o1** has been grasped and dropped into the bin, and therefore there is nothing more to do with it. The *pushing* action, still thinking in a deterministic scenario, is supposed to be done in a manner that the object will be moved far enough from the others in such a way that it can be considered isolated. Considering it is isolated from the others, if it was blocking some object to be pushed in a certain direction now it will not impede no more, and the same if an object impeded the current one, so the **block_dir_i** predicates are updated, and the **block_grasp** as well.

Geometric constraints The proposed planner until now does not include any geometric information, such as whether the robot can actually perform a certain action. It may happen that the object is outside the working space of the robot and this makes it hard to interact with. In such a case the robot will have a plan but it is impossible its execution, this make the plan unfeasible.

This aspect has to be taken into account and it can be done using a common technique in task planning such as **backtracking** [11]. The backtracking technique is about compensating the gap between planning with abstract actions and some physical and geometric constraints such as kinematics. Backtracking is based on updating the initial state used to get the current plan considering the geometric and kinematics constraints, and replan again with the updated initial state, and obtaining in this way a different plan that could be feasible. The obtained plan will still be evaluated and backtracked until a feasible plan is obtained or there exists no plan. The pseudo algorithm is shown in Algorithm 1.

This technique is a very common technique for planners that plans symbolically and have no informations about geometric constraints. This is considerable limitation of symbolic planning but makes it very fast to execute.

Let's consider to include the geometric information, such as path planning or inverse kinematics, into the geometric planner. The inverse kinematics by itself is an expensive operation, overall considering the high degree of freedom of the WAM arm. The most of the times the objects are considered to lie inside the working space of the robot. Computing the inverse of Kinematics for each possible action for each object would be very expensive, and by experimental result the majority of the

Algorithm 1 Planning procedure

```
procedure GETPLAN( $state_{init}, goal$ )                                ▷ Initial state and goal
    repeat
         $plan = \text{GETFASTDOWNWARDPLAN}(state_{init}, goal)$ 
        if ISPLANEMPTY() then return NULL
        end if
         $action, object = \text{GETFIRSTACTION}(plan)$  ▷ Action to execute and object
        of interest
         $success = \text{HASIKSOLUTION}(action, object)$ 
        if  $\neg success$  then
             $state_{init} = \text{UPDATEINITIALSTATE}(action, object)$  ▷ Update  $state_{init}$ 
            adding a predicate that avoid to execute  $action$  on  $object$ 
        end if
    until  $success$  return  $plan$  ▷ Return a geometric feasible plan or no plan at
    all
end procedure
```

times the inverse kinematics is feasible and the majority of the plans are feasible. Moreover, due to the probability of the next states the planner has to replan at each frame, and computing the predicates means computing the inverse kinematics as well. All this would make the planner process extremely time consuming when it is not needed.

To consider the geometric unfeasibility of an action a predicate for each action is added: `ik_unfeasible_dir1 ?o`, `ik_unfeasible_dir2 ?o`, `ik_unfeasible_dir3 ?o`, `ik_unfeasible_dir4 ?o` and `ik_unfeasible_grasp ?o`. The backtracking procedure will try to execute the action by first computing a solution for the inverse kinematic for the path to execute, the pushing action will be a sequence of points in the Cartesian space and the grasping action will be composed of a pre-grasping pose and a grasping pose, and then if that action on a particular object is not possible to execute one of those predicates is added.

It is possible that an object cannot be grasped on a certain pose but it can be move in a new pose in which it can be grasped. Therefore the pushing actions also will include the effect that the object of interest will have a solution in the inverse kinematic for the new pose. This is usually a rare case but it may happen. Of course, since the planner has no geometric information, it does not know in which pushing direction the object can be reach a position where it can be grasped, therefore this will be a totally random choice up to the planner. In the case the object is in a pose where it cannot be neither grasped or pushed because of the inverse kinematic, first the planner will return as a solution to grasp it, then the it will replan and the solution will be to push it in one direction and grasp it, and so until no action can be executed and there exist no solution for the plan. In the best case, the planner will chose to push the object in direction in which the next pose will be more likely a graspable one.

It may happen that the object outside the working space of the robot block the execution of the task because if the planner insists to return as first action an action with that object, then it will find that with that object is not possible to interact with and therefore there exist no solution on the planner procedure. In our experiments we observed that actually the planners tends to leave the "problematic" object as the last one. This behaviour is due to the manner it chooses the actions,

when there exists more feasible plans it chooses one of them (both optimal), this makes that "problematic" object could or not block the execution of all the task. How previously said, it was observed that the objects that are difficult to manipulate for the lack of solutions in the inverse kinematics are usually abandon and there are the last ones to be manipulated. (CHECK THIS AFFIRMATION)

INCLUDE SIMPLE EXAMPLE

PDDL For clarity purposes, the PDDL syntax of the described action is here reported. For the *grasping* action its PDDL syntax is shown in listing ??.

```

1 (:action grasp
  :parameters (?o - obj)
  :precondition (and
    ; grasp it if there are no objects on its top
    (not (exists (?x - obj)(on ?x ?o)))
    ; grasp it if there is no object that blocks it
    ; to be grasped
    (not (exists (?x - obj)(block_grasp ?x ?o))))
  :effect (and
    ; the object "o" is grasped
    (grasped ?o)
    ; if the object was on top of other ones now it is
    ; no more on top of them
    (forall (?x - obj)
      (when (on ?o ?x) (not (on ?o ?x))))
      ; the grasped objects no more blocks other objects
    ; to be pushed or grasped
    (forall (?x - obj)
      (when (block_grasp ?o ?x) (not (block_grasp ?o ?x))))
    (forall (?x - obj)
      (and
        (when (block_dir1 ?o ?x) (not (block_dir1 ?o ?x)))
        (when (block_dir2 ?o ?x) (not (block_dir2 ?o ?x)))
        (when (block_dir3 ?o ?x) (not (block_dir3 ?o ?x)))
        (when (block_dir4 ?o ?x) (not (block_dir4 ?o ?x))))))
  )

```

Listing 3.1: PDDL syntax of the grasping action

For the *pushing* action its PDDL syntax is shown in listing ??.

```

1 (:action push_dir1
  :parameters (?o - obj)
  :precondition (and
    ; push in direction 1 only if there are no
    ; objects that block it along that direction
    (not (exists (?x - obj)(block_dir1 ?x ?o)))
    ; push it if it has no objects on top of it
    ; and if it not on top of other ones
    (not (exists (?x - obj)(on ?x ?o)))
    (not (exists (?x - obj)(on ?o ?x))))
  :effect (forall (?x - obj)
    (and
      ; once pushed the object is no more blocked in any
      direction
      ; and it no more blocks other objects to be moved
      (when (block_dir1 ?o ?x) (not (block_dir1 ?o ?x)))
      (when (block_dir2 ?o ?x) (not (block_dir2 ?o ?x)))
      (when (block_dir3 ?o ?x) (not (block_dir3 ?o ?x)))
      (when (block_dir4 ?o ?x) (not (block_dir4 ?o ?x)))
      (when (block_dir1 ?x ?o) (not (block_dir1 ?x ?o)))
      (when (block_dir2 ?x ?o) (not (block_dir2 ?x ?o))))
  )

```

```

21      (when (block_dir3 ?x ?o) (not (block_dir3 ?x ?o)))
22      (when (block_dir4 ?x ?o) (not (block_dir4 ?x ?o)))
23      ; once pushed it can be grasped and it no more
24      ; blocks other objects to be grasped
25      (when (block_grasp ?x ?o) (not (block_grasp ?x ?o)))
26      (when (block_grasp ?o ?x) (not (block_grasp ?o ?x))))))

```

Listing 3.2: PDDL syntax of the pushing action along direction 1

With such predicates the planner has all the basic information about the location of the objects on the scene. The symbolic predicates try to simplify the problem translating the geometric properties of the objects in symbolic ones. It is important to point out again that the system is deterministic, meaning that all the actions are supposed to give the resultant state with a probability of 1. Clearly the more uncertainty is related to the pushing action, the direction selection, based on the principal axis, does not take into account reliably the geometry of the object. The trajectory will be unlikely the desired one but a similar one. In the case the result of the action is not the expected one the replanning strategy will resolve the uncertainty problem. In this manner we are forced to consider that some interactions between objects are allowed but tried to be avoided as much as possible.

The way of computing the predicates will be discussed in detail in the next sections.

3.4. Limitations

The planner here presented has some clear limitations, these are related to the lack of the probability and mainly to the lack of geometric information. As previously said, the pushing action is supposed to push the object far enough from the others ones, this will be true only if the pushing will be performed up to infinite, and this is not the case. The pushing is performed taking into account the geometry of the manipulated object, without taking care about the surrounding objects geometry and poses. This problem is though to be resolved by replanning, that is, if an object has been pushed but not enough, the planner will return a sequence where the first action is again to push that object. This is actually what the planner does, but since for the most of cases, pushing along directions 1 and 2 is feasible, for the planner both directions are feasible to resolve the problem. So what happens is that the planner could first return as solution to push a certain object along direction 1, and then, when replanning, to push the object along direction 2. What is here commented is a infinite loop. This happen because the planner has no information between consecutive frames, that is, it consider the problem as a separated one from the previous frame. This fact is a very undesirable one and there are several scenarios that can present a behaviour like that one. The promising part is the combination of the grasping and pushing actions. It has been observed that the majority of times, although for complex scenarios, the solution is mainly based on a proper sequence of grasping action, while the pushing action is an auxiliary action which is rarely used. Taking this into account, and the presented limitation, is very likely that once an object has been pushed, ones of the objects which cannot be grasped before can now be grasped, avoiding such undesired infinite loop.

4. Algorithm

In this chapter the algorithm is discussed, presenting how the objects are recognized, and how the symbolic predicates are obtained.

4.1. Object Localization

For the planner, to know how to move the object, is fundamental knowing the objects on the scenario, therefore they need to be detected. More correctly, they need to be segmented since the algorithm is dealing with unknown objects, it is not going to recognize an objects as a particular one, but it segments the objects.

This stage is composed on two steps:

1. Detecting the table top objects
2. Segmenting the table top objects

The depth sensor is recording a depth map of a table, the algorithm has therefore to detect first the table, and so the objects that stands on top of it, and then segmenting them. We don't want to segment the entire image, if so, the table will be segmented as an object and the floor as well.

4.1.1. Tabletop Object Detection

The strategy for the tabletop object detection phase is composed of 3 different steps:

1. **Table plane estimation** (by RANSAC): the points of the table are detected estimating first a plane in the point cloud, all the points which belong to such a plane are the points of the table.
2. **2D Convex Hull of the table**: having the points of the table a 2D convex hull is computed in order to get a 2D shape containing those points.
3. **Polygonal prism projection**: all the points are projected on the table plane previously estimated and all the points which projections belong to the 2D convex hull are considered to be points of tabletop objects.

The steps of this tabletop object detection algorithm are described in Figure 4.1 for the point cloud in Figure 4.1a.

4.1.2. Object Segmentation

Once the objects on the table are detected the following phase is to segment them in order to get a point cloud per object.

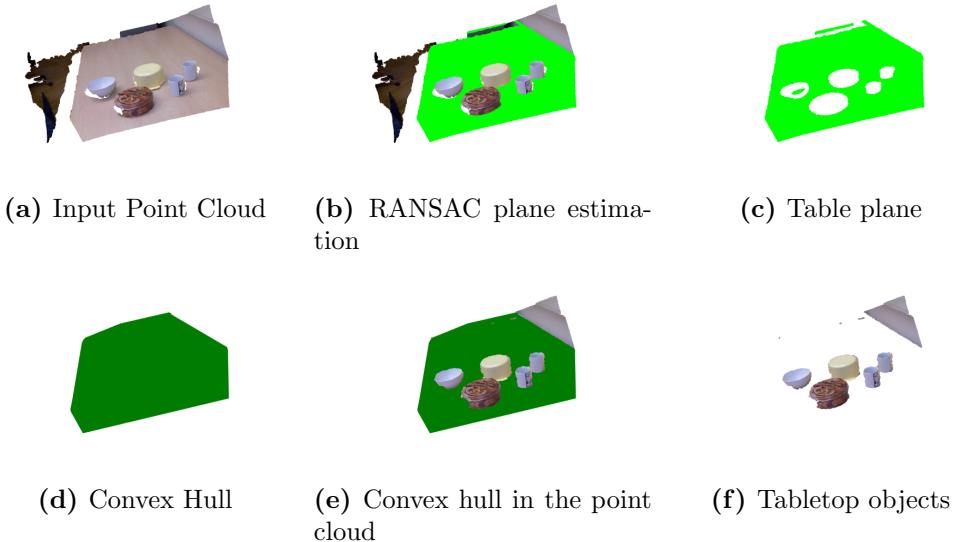


Figure 4.1: Object Segmentation: Given the point cloud (a), the estimated table’s plane is obtained (b and c), its convex hull is extracted (d and e), and the tabletop objects are obtained by a polygonal prism projection (f).

Supervoxel

For their segmentation the supervoxel concept is used. A supervoxel is a group of voxels that share similar characteristics.

In this work the supervoxels are computed with the Voxel Cloud Connectiviy Segmentation (VCCS) algorithm [36], which was proposed in 2014 and gives good improvements with respect to previous state of the art methods, and still it is able to be used in online applications.

The algorithm works in 4 main steps:

- Voxelizing the point cloud
- Creating an adjacency graph for the voxel-cloud
- Creating seeds for the initial supervoxels centres
- Each seed is described by 39 features that describe spatial coordinates, colors and local surface model properties. Then a distance metric based on these features is defined.
- Clustering the voxels into supervoxles iteratively by means of the distance metric, the adjacency graph, and the search volume of the supervoxel.
- Once the search of all supervoxel adjacency graphs has been concluded, the centres of each supervoxel is updated by taking the mean of all its constituents.

Local Convex Connected Patches Segmentation

Once the supervoxels of the objects are computed, they can be clustered in order to segment the objects. Papon et al. also proposed a segmentation algorithm based on their supervoxel technique, called *Local Convex Connected Patches Segmentation* (LCCP) [40]. This algorithm permits to segment objects by clustering together

adjacent convex supervoxels. The algorithm is quite simple but very good for segmentation of objects that have convex shapes, in Figure 4.2 the algorithm is briefly described.

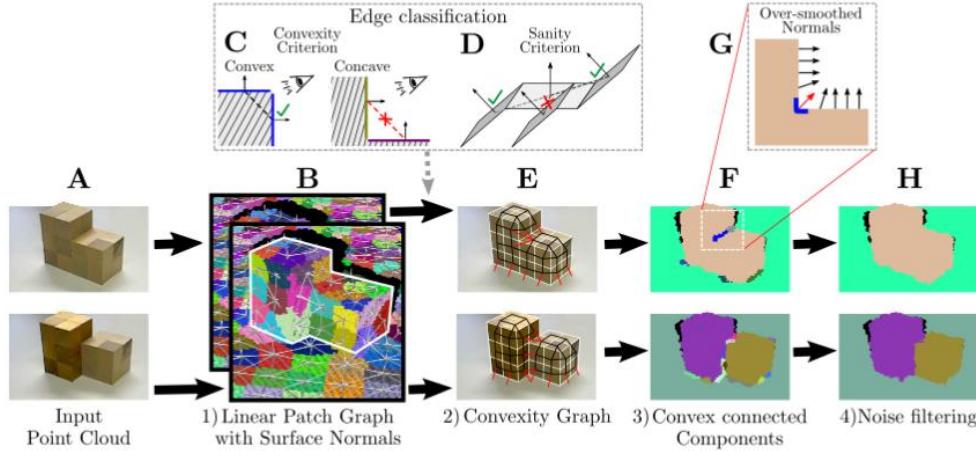


Figure 4.2: LCCP algorithm's structure. Reproduced from [40]

It clusters all the adjacent convex supervoxels (patches) using two criterion:

- Extended criterion: to consider two adjacent patches convex, both must have a connection to a patch which is convex with respect both patches
- Sanity Criterion: check if the adjacent patches which can be considered as convex present geometric discontinuities (see point D of Figure 4.2), in this case they are not considered as valid to form a cluster.

Then, due to the smoothed normals that could appear in some edges of the objects (point G Figure 4.2), the algorithm merges the clusters that are composed of few supervoxels to the biggest adjacent cluster.

By tuning properly the parameters of the segmentation algorithm the objects can be correctly segmented obtaining for one of them a point cloud.

ADD EXAMPLE IMAGES

4.2. Predicates Computations

Once the objects are segmented, we will have one point cloud per object, next we have to translate the scenario composition into symbolic predicates.

In Chapter 3 the predicates that are used are described, in this section their computation is presented.

Before going in detail on their computation some general concepts are first presented.

4.2.1. Tools

- collision detection
- pca
- table projection - convex hulls

To understand if an object impeded a certain action, such as the pushing along a direction, we have to check if along the desired trajectory the pushed object will collide with a certain one. The collision detection is therefore a crucial step for the predicates computation. For the collision checking the Flexible Collision Library (FCL) [35] has been used. This library allows to define the collision problem in a simpler manner than other more famous collision libraries such as *Bullet* [8], and it can work with different objects shapes such as box, spheres, cone, convex, mesh and octree. The main library used in this work is the Point Cloud Library (PCL) [6], which allows some methods to create an object shape from a point cloud. The mesh shape has been thought to use since it can be easily obtained by using the *convex hull* algorithm.

Explain convex hull.

The PCL convex hull algorithm returns the convex hull as a triangular mesh, which can be used by the FCL for the collision detection.

The pose of the Kinect camera is above the table pointing forward it. In such a pose the Kinect is only able to see mainly the top side of the objects, and not their sides. — discussed the convex hull relating into to the necessity for the collision detection.

Somewhere talk about the limitation of the segmentation that is we assume it to be perfect because if not the planner will return likely an infeasible plan

4.2.2. Predicate: block_grasp

4.2.3. Predicate: on

4.2.4. Predicate: block_dir_i

5. Implementation

Talk about ROS and real robot.

6. Experiments

In this chapter the performance of the proposed task planner for the considered task is evaluated by performing several experiments with the aim to study the quality, its benefits and its limitations.

To get the quality of the task planner the strategy defined in the article [10] by Amignoni et al. In that article they proposed a general method to evaluate the quality of a task planner. First they define a functionality benchmark¹ thought 4 elements:

- Description: a high level, general, description of the functionality
- Input/Output: the information available to the module implementing the functionality when executed, and the expected outcome.
- Benchmarking data: the data needed to perform the evaluation of the performance of the functional module.
- Metrics: algorithms to process benchmarking data in an objective way.

Regarding the benchmarking of a task planner those elements are defined as:

Description: Given a initial state and a goal state find a set of actions to reach the goal.

Input/Output: The inputs are the PDDL domain and problem descriptions. The output is a plan expressed as sequence of actions.

Benchmarking data: Some data has to be supplied by the robot, including: the input given to the planner, the output (i.e., the plan) delivered by the planner, and the performance data for the planner (e.g., memory and time needed). According to the planning approach, other data might be required. Other data will be possibly collected when a plan is actually executed, including: time needed to execute each action and time needed to execute the whole plan.

Metrics: Metrics are mainly related to good plan quality and to solving time. The target of benchmarking task planning is twofold: to asses the quality of the plan produced (e.g., in terms of the cost for executing it and its likelihood to succeed) and to assess the quality of the planning process (i.e., how fast is the plan generated and how much resources are needed to do so). Possible metrics are:

¹A functionality benchmark is a benchmark of specific robotic task. In the case of this thesis a benchmark of the task planner.

- Levenshtein distance² is a metric that specifies the distance between two strings, in this case a plan can be thought as a sequence of character, where each action with a determined object of interest is a character of the string. The Levenshtein distance between the two lists is calculated to measure the quality of the plan. Levenshtein distance $L_{A,B}(|A|, |B|)$ between two lists A and B is calculated as

$$\begin{aligned} L_{A,B}(i, j) &= \max(i, j) \quad if \quad \min(i, j) = 0 \\ L_{A,B}(i, j) &= \min(L_{A,B}(i - 1, j) + 1, \\ &\quad L_{A,B}(i, j - 1) + 1, \\ &\quad L_{A,B}(i - 1, j - 1) + 1_{a_i \neq b_i}) \quad otherwise \end{aligned}$$

where $|A|$ is the length of plan A, and $1_{a_i \neq b_i}$ is equal to 0 if the action a_i is the same of action b_i . This metric will be used to evaluate the distance between the optimal plan returned by the planner in the first frame and the real executed plan.

- Number of actions correctly performed in the correct order (e.g., number of objects delivered to their correct destination).
- Time required to construct a plan and time to perform it.

We will also evaluate the difference of executing the initial plan and executing the plan with the replanning technique in terms of time and number of correct actions. As well also the number of bad, or dangerous, object manipulations due to the lack of replanning, lack of probability and geometric constraints in the planner .

It is important highlighting that what has been done is not a benchmarking since no comparison is done with other state of the art planners, but the commented guideline for benchmarks has been used in order to get some useful metrics in order to assert the quality of the planner.

Kinds of possible experiment to do:

- show experiments first trying only separating the object (infinite loop behaviour) and then coupling the grasping action.
- Show the results of the planner with the pushing length chosen not accordingly to the surrounding objects but accordingly to the pushed one. From that comment the improvements.
- Show normal experiment where everything works
- Show situation why we are going to consider the risk as cost in the action.

Metrics to consider:

- number of objects
- segmentation time
- Total solution time
- planning time

²https://en.wikipedia.org/wiki/Levenshtein_distance

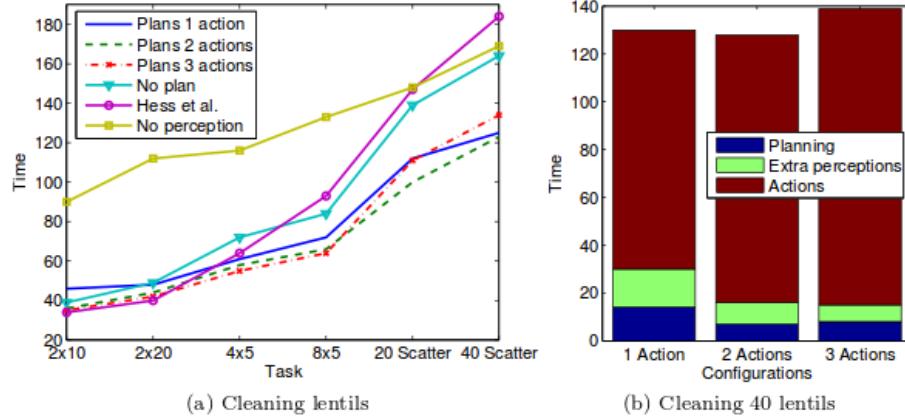


Figure 8: (a) Skill of moving lentils to a container. The results shown are the mean over 5 runs. (b) Time for the case of moving 40 scattered lentils to a container. Extra perceptions include the movements and perceptions needed to tackle occlusions by taking the arm out of the field of view (see Section 4.2.3).

- total solution
- Inverse kinematics solution time (this will probably used to justify the back-tracking method)
- Predicates time computation (maybe add collision checking time in)

Find a way to get the computational complexity.

7. Further Work

Bibliography

- [1] Description of the shop project <https://www.cs.umd.edu/projects/shop/description.html>.
- [2] Gourmand Home Page.
- [3] PROST Home Page.
- [4] SHOP Home Page.
- [5] Fast Downward Home Page.
- [6] <http://www.openperception.org/> - open perception.
- [7] PDDL - The Planning Domain Definition Language. Technical report, CVC TR-98-003/DCS TR-1165, Yale Center for Computational Vision and Control, 1998.
- [8] Bullet <http://wiki.ros.org/bullet>, 2015.
- [9] Nichola Abdo, Henrik Kretzschmar, Luciano Spinello, and Cyrill Stachniss. Learning manipulation actions from a few demonstrations. In *Robotics and Automation (ICRA), 2013 IEEE International Conference on*, pages 1268–1275. IEEE, 2013.
- [10] Francesco Amigoni, Andrea Bonarini, Giulio Fontana, Matteo Matteucci, Viola Schiaffonati, Aamir Ahmad, Iman Awaad, Jakob Berghofer, Rainer Bischoff, et al. General evaluation criteria, modules and metrics for benchmarking through competitions. 2014.
- [11] Julien Bidot, Lars Karlsson, Fabien Lagriffoul, and Alessandro Saffiotti. Geometric backtracking for combined task and motion planning in robotic systems. *Artificial Intelligence*, pages –, 2015.
- [12] Stéphane Cambon, Fabien Gravot, and Rachid Alami. A robot task planner that merges symbolic and geometric reasoning. In Ramon López de Mántaras and Lorenza Saitta, editors, *ECAI*, pages 895–899. IOS Press, 2004.
- [13] Rui Coelho and Alexandre Bernardino. Planning push and grasp actions: Experiments on the icub robot.
- [14] Hao Dang and Peter K. Allen. Semantic grasping: Planning robotic grasps functionally suitable for an object manipulation task. In *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS 2012)*, pages 1311–1317. IEEE, October 2012.

- [15] Richard Dearden and Chris Burbridge. An approach for efficient planning of robotic manipulation tasks. In *ICAPS*. Citeseer, 2013.
- [16] Richard Dearden and Chris Burbridge. Manipulation planning using learned symbolic state abstractions. *Robotics and Autonomous Systems*, 62(3):355 – 365, 2014. Advances in Autonomous Robotics — Selected extended papers of the joint 2012 {TAROS} Conference and the {FIRA} RoboWorld Congress, Bristol, {UK}.
- [17] Mehmet Dogar and Siddhartha Srinivasa. Push-grasping with dexterous hands: Mechanics and a method. In *Proceedings of 2010 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS 2010)*, October 2010.
- [18] Mehmet Dogar and Siddhartha Srinivasa. A framework for push-grasping in clutter. In Nick Roy Hugh Durrant-Whyte and Pieter Abbeel, editors, *Robotics: Science and Systems VII*. MIT Press, July 2011.
- [19] Mustafa Ersen, Melodi Deniz Ozturk, Mehmet Biberci, Sanem Sariel, and Hulya Yalcin. Scene interpretation for lifelong robot learning. In *Proceedings of the 9th international workshop on cognitive robotics (CogRob 2014) held in conjunction with ECAI-2014*, 2014.
- [20] Mustafa Ersen, Sanem Sariel Talay, and Hulya Yalcin. Extracting spatial relations among objects for failure detection. In *KIK@ KI*, pages 13–20, 2013.
- [21] Keliang He, Morteza Lahijanian, Lydia E. Kavraki, and Moshe Y. Vardi. Towards manipulation planning with temporal logic specifications. In *2015 IEEE Intl. Conf. Robotics and Automation (ICRA)*, pages 346–352, Seattle, WA, May 2015. IEEE.
- [22] Malte Helmert. The fast downward planning system. *J. Artif. Intell. Res.(JAIR)*, 26:191–246, 2006.
- [23] Tucker Hermans, James M. Rehg, and Aaron F. Bobick. Guided pushing for object singulation. In *IROS*, pages 4783–4790. IEEE, 2012.
- [24] Dov Katz. *Interactive perception of articulated objects for autonomous manipulation*. University of Massachusetts Amherst, 2011.
- [25] Dov Katz, Moslem Kazemi, J. Andrew (Drew) Bagnell, and Anthony (Tony) Stentz . Clearing a pile of unknown objects using interactive perception. In *Proceedings of IEEE International Conference on Robotics and Automation*, March 2013.
- [26] Dov Katz, Arun Venkatraman, Moslem Kazemi, J Andrew Bagnell, and Anthony Stentz. Perceiving, learning, and exploiting object affordances for autonomous pile manipulation. *Autonomous Robots*, 37(4):369–382, 2014.
- [27] Steven M. LaValle. *Planning Algorithms*. Cambridge University Press, May 2006.
- [28] Iain Little, Sylvie Thiebaux, et al. Probabilistic planning vs. replanning. In *ICAPS Workshop on IPC: Past, Present and Future*, 2007.

- [29] Tomás Lozano-Pérez and Leslie Pack Kaelbling. A constraint-based method for solving sequential manipulation planning problems. In *Intelligent Robots and Systems (IROS 2014), 2014 IEEE/RSJ International Conference on*, pages 3684–3691. IEEE, 2014.
- [30] Kevin M Lynch and Matthew T Mason. Stable pushing: Mechanics, controllability, and planning. *The International Journal of Robotics Research*, 15(6):533–556, 1996.
- [31] Bhaskara Marthi, Stuart J Russell, and Jason Wolfe. Angelic semantics for high-level actions. In *ICAPS*, pages 232–239, 2007.
- [32] Tekin A Mericli, Manuela Veloso, and H Levent Akin. Achievable push-manipulation for complex passive mobile objects using past experience. In *Proceedings of the 2013 international conference on Autonomous agents and multi-agent systems*, pages 71–78. International Foundation for Autonomous Agents and Multiagent Systems, 2013.
- [33] Lorenz Mösenlechner and Michael Beetz. Using physics- and sensor-based simulation for high-fidelity temporal projection of realistic robot behavior. In *AIPS*, 2009.
- [34] Melodi Ozturk, Mustafa Ersen, Melis Kapotoglu, Cagatay Koc, Sanem Sariel-Talay, and Hulya Yalcin. Scene interpretation for self-aware cognitive robots. 2014.
- [35] Jia Pan, Sachin Chitta, and Dinesh Manocha. Fcl: A general purpose library for collision and proximity queries. In *Robotics and Automation (ICRA), 2012 IEEE International Conference on*, pages 3859–3866. IEEE, 2012.
- [36] Jeremie Papon, Alexey Abramov, Markus Schoeler, and Florentin Wörgötter. Voxel cloud connectivity segmentation - supervoxels for point clouds. In *Computer Vision and Pattern Recognition (CVPR), 2013 IEEE Conference on*, Portland, Oregon, June 22-27 2013.
- [37] R. Reiter. A logic for default reasoning. In M. L. Ginsberg, editor, *Readings in Nonmonotonic Reasoning*, pages 68–93. Kaufmann, Los Altos, CA, 1987.
- [38] Stuart J. Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Pearson Education, 2 edition, 2003.
- [39] T. Siméon, J.-P. Laumond, and C. Nissoux. Visibility-based probabilistic roadmaps for motion planning. *Advanced Robotics*, 14(6):477–493, 2000.
- [40] S. C. Stein, M. Schoeler, J. Papon, and F. Woergoetter. Object partitioning using local convexity. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR) 2014*.
- [41] Nikolaus Vahrenkamp, Tamim Asfour, and Rudiger Dillmann. Robot placement based on reachability inversion. In *Robotics and Automation (ICRA), 2013 IEEE International Conference on*, pages 1970–1975. IEEE, 2013.
- [42] Anna Yershova and Steven M. LaValle. Deterministic sampling methods for spheres and SO(3). In *Proceedings of the 2004 IEEE International Conference on Robotics and Automation, ICRA 2004, April 26 - May 1, 2004, New Orleans, LA, USA*, pages 3974–3980, 2004.