

UNIVERSITAT POLITÈCNICA DE CATALUNYA

MASTER THESIS

MASTER'S DEGREE IN AUTOMATIC CONTROL AND ROBOTICS

Task planning for table clearing of cluttered objects

Author:
Nicola COVALLERO

Director:
Guillem ALENYÀ RIBAS
Co-director:
David MARTÍNEZ MARTÍNEZ
Tutor:
Guillem ALENYÀ RIBAS

June 2016

Abstract

Manipulation planning is a field of study with increasing interest, it merges manipulation skills with an artificial intelligence system that is able to find the optimal sequence of actions in order to solve manipulation problems. Those problems are complex since they involve a mixture of symbolic planning and geometric planning.

In this thesis we present a planning system for clearing a table of cluttered objects, which completely reasons at a symbolic level postponing the reasoning with geometric constraints through backtracking. The task we face is a complex task that should include geometric information, such as the objects positions, during the planning step. We show that the planning system developed is able to handle efficiently the task with symbolic predicates.

The main contribution of this thesis is a planning system able to solve a wider variety of scenarios for clearing a table of cluttered objects. Grasping actions alone are not enough, sometimes pushing actions are needed in order to move an object in a pose in which it can be grabbed. The planner here presented is able to understand when, and where, to push an object and why it is not possible to grasp it.

This work shows how some geometric problems can be efficiently handled by reasoning at an abstract level through symbolic predicates when such predicates are chosen correctly. The advantage of this kind of planning system is a reduction in computation time and also an easier implementation.

Contents

Contents	2
1 Introduction	3
1.1 Problem Approach	3
1.2 Set Up	4
2 State of the art in manipulation planning	7
3 Planning system	9
3.1 Task Planners Review	9
3.2 Planner	9
3.3 Symbolic Predicates	10
4 Implementation	17
4.1 Object Localization	17
4.1.1 Tabletop Object Detection	17
4.1.2 Object Segmentation	18
4.2 Background	20
4.3 Actions Execution	24
4.3.1 Pushing	24
4.3.2 Grasping	27
4.4 Predicates Computations	28
4.4.1 Predicate: <code>block_grasp</code>	28
4.4.2 Predicate: <code>on</code>	28
4.4.3 Predicate: <code>block_dir_i</code>	30
5 Software design	33
6 Experiments	34
7 Conclusions	37
7.1 Conclusions	37
7.2 Limitations	37
7.3 Future Work	37
Bibliography	38

1. Introduction

Robotic manipulation of objects is an increasing field of research which has struggled researchers from many years ago. In several industrial environments robots can be easily programmed considering the objects to be known a priori, i.e. the manipulation is always the same, and usually robot operations avoid to deal in cluttered scene, but this situation needs a workspace designed in a manner to provide to the robot a non cluttered scene. But there are situations in which robots with enhanced intelligence can be useful. An example in which a robot could face a cluttered scenario is the one of the *Amazon Picking Challenge* [10], which has the objective to provide a challenge problem to the robotics research community that involves integrating the state of the art in object perception, motion planning, grasp planning, and task planning to manipulate real-world items in industrial settings such ones human operators face in Amazon's warehouse. Joey Durham from Amazon Robotics describes the challenges of this competition with the following sentence:

“A robot for picking purposes must possess a lot of skills: The selected item must be identified, handled and deposited correctly and safely. This requires a certain level of visual, tactile and haptic perceptive skills and great handling capabilities.”

In this thesis we are going to investigate a simple approach for a complex manipulation problem. Our approach tries to replicate the human reasoning during the manipulation of cluttered objects in table clearing tasks. The objective is removing all the objects onto a table considering that their poses could limit the manipulation actions.

This thesis has been developed in the **Institut de Robòtica i Informàtica Industrial** (IRI) in the Perception and Manipulation laboratory with the supervision of *Guillem Alenyá Ribas* as director and *David Martínez Martínez* as co-director.

This work is structured as follow. First an introduction to the problem we want to solve is presented in the Chapter 1. Then in Chapter 2 a review of the current state of the art in manipulation planning is done. The planner developed will be commented in Chapter ?? and the implemented algorithm to compute the predicates and control the robot will be commented in Chapter ???. In Chapter 4 the softwares employed for the implementation of the code will be presented as well the algorithm's structure. Finally some experiments and conclusions about this work are discussed in Chapter 6 and 7 respectively.

In this chapter first an introduction to the problem and the approach adopted is commented, then the experimental set up is presented since it will be useful to understand some techniques and choices for the algorithm which will be discussed in next chapters.

1.1. Problem Approach

In this section the approach to solve the planning problem is described. The strategy to solve the problem is inspired to a human-like solution. A human would solve such task with three main actions: grasp an object, push it or drag it. When it is not possible to grasp an

object, because other objects hinder the human to put the hands in the proper way to grasp the desired one, he/she has to interact with the scene in order to put the desired object in a pose where it can be grasped.

Dragging is a very complex action which will require the robot to lay the end effector on the top of an object. This is a very hard action which would imply an implementation of a reliable controller, moreover the goal is not only to move the object but also not to ruin it.

Pushing action is an action easier than dragging since it only requires to put the end effector in a certain pose and then push the object by moving the end effector. However the pushing action is complex since to push an object making it following a certain path the pushing pose has to be chosen taking into account the shape of the object and a controller would be needed in order to correct the pose along the path. Despite this we handled to perform that action in a easy way by considering only objects with basic shapes.

For these considerations the main actions the robot has to use in order to interact with the objects are grasping and pushing. Grasping is the most important action since it lets to take an object from the pile of objects and drop it somewhere, for instance into a bin, clearing in this way the table. There exist different works facing the same task by focusing only in grasping [17] [37]. The pushing becomes useful considering the problem that two adjacent objects could not be grasped if they are so close such that the robot's gripper, when attempting to grasp an object, is going to collide with the adjacent one, making the object of interest ungraspable. From this consideration is necessary the pushing action, in order to separate adjacent objects which could mutually exclude themselves to be grasped.

Being observant to the philosophy of a human-inspired solution, the manipulation of the objects is intended in manner to avoid to interact with more than one object per action. That is, when we push, or grasp, an object only that object will change its pose.

The robot's intelligence is enhanced by a planning system which will return a sequence of actions in order to achieve the goal to clear the table. The robot will reason totally in an abstraction level by considering only symbolic predicates which will be obtained by hand-built mapping functions making simpler the reasoning step. It is important noting that this problem should includes important geometric information which are very useful for the planner, such as how much to push an object, in what direction and the inverse kinematic of the robot. This geometric information are translated into symbolic ones relaxing the planning phase, although these are not able to catch all the geometric information. Then through backtracking the robot will evaluate if the returned plan can be actually executed by the robot (i.e. evaluate if the action to execute involves the robot to move inside its working space), if not it updates the predicates and replans. We assume the world is perfectly known and the actions are not reliable. To overcome this the planner replans after the execution of each action. To detect the objects the Microsoft Kinect sensor is used as vision system and a geometry based segmentation is used to segment the objects in the scene.

1.2. Set Up

In order to make easier the understanding of the approach used, the set up of the environment the robot will work in is here presented.

The robot used is a Barret WAM arm, which is a 7 degree of freedom (DoF) manipulator arm shown in Figure 1.1a. The WAM is noteworthy because it does not use any gears for manipulating the joints, but cable drives, so there are no backlash problems and it is both fast and stiff. Cable drives permit low friction and ripple-free torque transmission from the actuator to the joints. To detect the objects a Kinect camera, a RGB-D sensor, is employed (Figure 1.1b).



(a) Barrett WAM arm



(b) Microsoft Kinect sensor

Figure 1.1: Robot and vision sensor.

To manipulate the objects the manipulation skills of the robot are enhanced by a gripper designed in the IRI institute and actuated by Dynamixel motors. Such a gripper is depicted in Figure 1.2 from several point of views. Its closing width¹ is 3 centimetres while the opening width² is of 8 centimetres, therefore we are constrained to grasp objects with a width in the range [3 ÷ 8]cm.

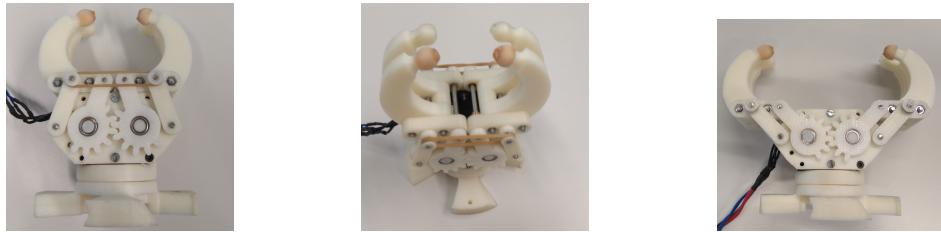


Figure 1.2: Gripper used for the experiments

For the task planner, as the reader will see in Chapters ?? and ??, the model of the gripper will be an important resource in order to compute the predicates.

The gripper will be modelled measuring some principal elements such as: finger's width and height, gripper's width, height and deep, closing and opening width. The modelling procedure is depicted in Figure 1.4. The resulting model is a simple triangle mesh which includes all the important geometric information of the gripper. Such a simple model allows



Figure 1.3: Gripper and WAM.

the collision algorithm commented in Chapter ?? to check for collision in just few millisecond. A more detailed and complex model would only lead as benefit a higher precision, but for this kind of task an extremely accurate precision is not needed, and it will slow down the algorithm because the collision checking procedure would be more computationally ex-

¹Distance between the fingers when the gripper is closed.

²Distance between the fingers when the gripper is open.

pensive. The gripper will be mounted in the end effector of the robot as shown in Figure 1.3.

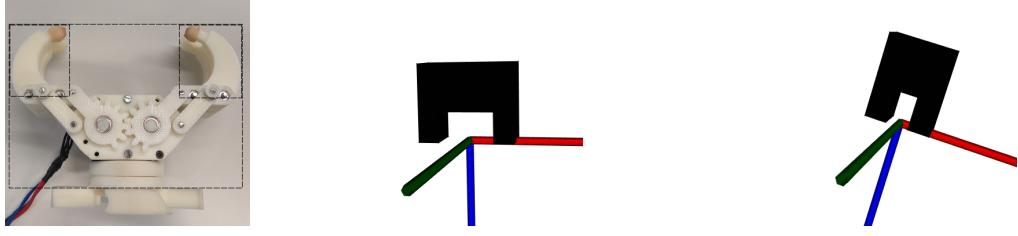


Figure 1.4: At the left the principal elements measured are highlighted for the opened gripper model. The gripper mesh model is here shown in the PCL visualizer. The red, green and blue axis are respectively the x, y and z axis.

The scenario the robot is going to work in is composed of a table and the objects will lay on top of it. In Figure 1.5a the main elements of the set up are highlighted. The WAM arm's base is in a fixed position with respect the table and the Kinect camera is located on top of the table pointing down. The Kinect is calibrated and the fixed transformation between the Kinect's frames and the base frame of the robot is known, so all the points measured by the Kinect can be expressed in coordinates with respect the robot's base frame. An example of a cluttered scene the robot is going to deal with in this thesis is depicted in Figure 1.5b, and the same scene seen by the Kinect is shown in Figure 1.5c. With respect the Kinect's view the robot is located at the top side of the image, outside the field of view of the Kinect, in order to avoid to detect the arm as an object. If the robot arm would be present in the depth image we should apply some algorithms to identify the robot arm, which could also produce important occlusions, that is hide some objects from the Kinect's view. In order to avoid all this, the depth images considered are the ones provided by the Kinect when the robot is no more in the Kinect's view (these poses are known a priori).

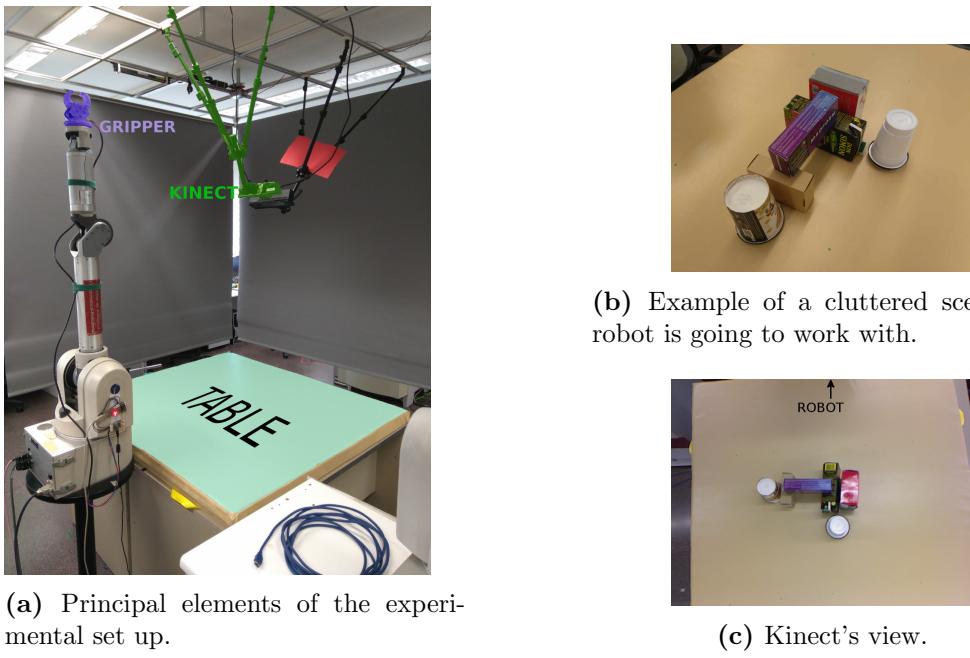


Figure 1.5: Experimental set up, and an example of a cluttered scene the robot is going to interact with.

2. State of the art in manipulation planning

Many manipulation planning approaches (see [26] for an overview) assume that the task can be treated as a geometric problem, with the goal to place the objects in their desired positions. Planning is essentially done with a mixture of symbolic and geometric states, this requires a set of symbolic predicates that correspond to geometric relationships. Cause generating geometric paths is very time consuming, these hybrid planners could be too slow to be used in real applications.

Dogar and Srinivasa [14] proposed a framework for planning in cluttered scenes using a library of actions inspired by human strategies. They designed a new planner that decides which objects to move, the order to move them, where to move them, and it chooses the manipulation actions to use on these objects, and it accounts for the uncertainty in the environment all through this process. The planning system first attempts to grasp the goal object, than if it is not possible it identifies what is the object that obstacles the action, then such an object is added to a list of the objects that have to be moved. They move those objects in whatever position such that makes the goal feasible. Their work is the most similar to our, but their planning system cannot be directly applied to a table clearing task. The goal is a single object at a time, then to grasp another object they need to replan. The planning system proposed in this thesis is able to perform more complex plans.

A recent alternative proposed by Mösenlechner and Beetz [30] is to specify goals symbolically but evaluate the plan geometrically. The idea is to use a high-fidelity physics simulation to predict the effects of actions and a hand-built mapping from geometric to symbolic states. Planning is conducted by a forward search, and the effects of actions determined by simulating them, and then using the mapping to update the symbolic state. Their method is promising, by the way the integration of a simulator into the planning system is not straightforward since it requires to know the physic of the manipulated objects. Moreover the authors didn't tested their planning system with a complex scene like the one faced in this thesis. Our planning system doesn't use any simulator, instead it uses a simple prediction algorithm to understand how the objects should be manipulated, leading to an easy implementation and a fast execution in the planning step.

In [9] the authors address a problem similar to the one of this thesis. The authors coupled together pushing and grasping actions for a table manipulation task. They use the concept of reachability [38] for each kind of action to quickly exclude impossible poses of the gripper of the planning stage, creating a reliable plan suitable for real-time operation. The authors model the planning problem through MDP (Markov Decision Process), discrediting the world in grid cells and assigning to each one a push and grasp vector defined by the reachability concept. Unfortunately the authors were not able to test the grasp skill with their planning system. Our approach skips the modelling of the probability by considering to replan after each action execution, this eases our planning phase. Moreover, while their method needs to be adapted to each robot, to built a reachability map, our method can be

directly integrated in any robotic manipulator.

Machine learning techniques also have been applied to task planning. To plan complex manipulation tasks, robots need to reason on a high level. Symbolic planning, however, requires knowledge about the preconditions and effects of the individual actions. In [5] the authors proposed a practical approach to learn manipulation skills, including preconditions and effects, based on teacher demonstrations. With few demonstration the authors proposed a method to teach the robot the preconditions and effects of individual actions. Their method furthermore enables the robot to combine the learned actions by means of planning to solve new tasks that are more complicated than the learned individual actions. Although this work looks promising the problem of this approach is that it moves the problem from identifying the right predicates to describe the problem to identifying the right features that let a correct learning of the predicates.

In [13] Dearden and Burbridge proposed an approach for planning robotic manipulation tasks which uses a learned mapping between geometric states and logical predicates. The planner first plans symbolically, then applies the mapping to generate geometric positions that are then sent to a path planner. They try to fit a probability distribution function to the geometric states to map them to a symbolic state. This work sounds promising but the cluttered scene this thesis is going to treat is quite complex to be treated with such a method. Moreover the same authors applied that technique [12] for a simple task which involves putting some cups (singulated between themselves) onto a tray. Although this approach does not involve any hand built mapping function and the results are good, their approach is very time consuming for real applications. The hand-built function still have a good benefit in terms of computational cost, although they need to be adapted to the problem. Despite this also learning the predicates requires to be adapted to the problem.

Compared to the state of the art this thesis proposes a planning system for table clearing tasks of cluttered objects, it plans completely at a symbolic level, efficiently and it is low time consuming (the time to get a plan is usually less than 0.5 seconds). The fast computation of the plan is able to make this system suitable for replanning, overcoming in this way the lack of probability and geometric information. Moreover no other works have been found that face the table clearing of cluttered objects problem, such as the one in Figure 1.5b. It will be shown that the lack of geometric constraints includes some limitations to the system but it plans still efficiently.

3. Planning system

In this chapter the general framework adopted is discussed, proposing a suitable task planning system. After the review of the current state of the art of task planners, a proper planner is chosen and then a suitable description to the table clearing problem is discussed.

3.1. Task Planners Review

To choose the proper planner for the task we evaluated three main categories of planners:

1. classical planners
2. hierarchical planners
3. probabilistic planners

Classical planners are characterized by environments which are fully observable, deterministic, finite and static (changes happen only when the agent acts) and discrete (in time, actions, objects...) [35]. A very well known classic planner is the Fast Downward planner [18].

Hierarchical planning, also called *Hierarchical Task Network*(HTN), works in a similar way to how it is believed that human planning works [28]. It is based on a reduction of the problem. The planner recursively decomposes tasks into subtasks, stopping when it reaches primitive tasks that can be performed directly by planning operators. This kind of planner needs to have a set of methods, where each method is a schema for decomposing a particular kind of task into a set of subtasks. For this kind of planning technique a well known planner is SHOP [31].

Probabilistic planning is a planning technique which considers that the environment is not deterministic but probabilistic. So the actions have a certain probability to obtain a certain state, and given an initial and final state the planner finds the solution path with the highest reward, which depends also on the probability. A well known example on which this kind of planners are build on is the Markov Decision Process. In this category two probabilistic planners that perform good in planning competitions are Gourmand [25] and PROST [24].

3.2. Planner

The problem involves a big amount of uncertainty due to the interaction of the robot with the environment. When the robot will interact with the objects it is very hard to predict correctly the position of the object after the execution, that is the next state, this is a crucial problem which should be considered. With a probabilistic planner, the planner will take into account what object has been moved and it will update the state obtaining a set

of states, each one with a certain probability. The probability also has to be modelled and it is highly depending on the form of the object.

Martinez et al. [29] faced the problem of cleaning a surface entirely by probabilistic symbolic planning. The problem they wanted to solve was characterized by a strong probability factor and a wrong effect of the action would have lead to replan and therefore to slow down the system. Another way to face the problem of probability is replanning after each executed action **cite Tomas Lozano**, or whenever the current state deviates from the expected one, generating a new trajectory from the current state to the goal. The plan's actions are considered deterministic, and the only useful action is actually the first one of the plan, after its execution the system replans again. Little et al. discussed in [27] the problem of when is more useful the probabilistic planning with respect a simple replanning system. They defined a planning problem *probabilistic interesting* if dead ends can be avoided or when there exist more feasible solutions with low probability such that coupled have a greater probability to reach the goal than another single trajectory with an higher probability.

They assert that unless a probabilistic planning problem satisfies all of the conditions to be *probabilistic interesting* then it is inevitable that a well-written replanner will outperform a well-written probabilistic planner. Moreover the authors do not negate the possibility that a deterministic replanner could perform optimally even for probabilistically interesting planning problems.

To conclude, a replanner makes more probabilistic problems practically solvable. In the other hand it suffers to be less robust than a probabilistic one.

Taking into account such considerations and that, except for rare cases, our planning problem is not *probabilistic interesting*, the problem has been thought to be solved thanks to a deterministic replanner.

A hierarchical planner would be a good choice if the problem presents some hierarchies, for instance in the case the goal is to clear several tables. Since the problem is about cleaning a single table the problem can be easily modelled through a classic planner and it is straightforward.

The planner chosen to develop the work is the **Fast Downward** planner [18], a very well know classic one. This planner is feature-wise complete, stable and fast in solving planning problem.

3.3. Symbolic Predicates

Once the planner has been chosen, the planning problem has to be formulated accordingly to the capabilities of the planner. Since we chosen a classic one, the problem is deterministic and it is generally formulated as a 6-tuple $\Pi = \langle S^d, s_o^d, G^d, A^d, T^d, c^d \rangle$ [27], where:

- S^d is a finite set of states;
- $s_o^d \in S^d$ is an initial state;
- $G^d \in S^d$ is a goal state;
- $A^d(s)$ is a set of applicable actions for each $s \in S^d$;
- $T^d(a, s) \in S^d$ is a deterministic transition function for all actions $a \in A^d(s)$ and states $s \in S^d$;
- $c^d(a)$ is the cost to apply action a .

The plans, or trajectories, τ_i are sequences of actions applicable from the initial state until the goal state. The cost of a trajectory $C(\tau_i)$ is the sum of the cost of the actions of the trajectory $C(\tau_i) = \sum_{a \in \tau} c^d(a)$. The optimal solution is the solution with less cost: $\tau^* = \min_{\tau_i} c(\tau_i)$. In this work we did not specify any cost to the actions, so the planner will return as solution the one with less number of actions.

The *Fast Downward* planner needs the problem to be formulated in the common format of PDDL (*Problem Domain Description Language*) [3]. In this section the symbolic predicates that have been considered in order to solve the problem are described.

The task this thesis faces is a common task done by humans, who think in order to find a feasible sequence of actions. Such a sequence is normally composed of actions that avoid the collision between the manipulated objects and the other ones, whenever possible. To do this we, as humans, think on what is going to happen if we manipulate an object in a certain way. The design of the problem has been inspired by such a reasoning way and therefore we will handle carefully the objects avoiding collisions between them. The reasoning part is entrusted to the planner, which needs a coherent problem description through symbolic predicates.

As described in the introduction, the system will be able to perform two types of actions: **pushing** and **grasping**. Grasping action is a necessary action in order to grasp an object and drop it into a bin, while the pushing action is an auxiliary action which has the aim to move an object in a pose that do not obstacle the solution of the problem. The pushing action is said to be auxiliary because it is not strictly necessary to solve every kind of problem, depending on the cluttered scene the grasping action could be enough. The blending of these two actions makes wider the range of problem this planner can solve.

The symbolic predicates are designed accordingly to the available actions trying to answer the following questions:

- When can an object be grasped?
- When can a object be pushed? In which direction?

Answering these questions the following predicates are defined:

- **grasped:** (`grasped o1`) means that object `o1` has been grasped and dropped into the bin. The goal is reached when all the object have been grasped.
- **on:** (`on o1 o2`) means that object `o1` stands on top of object `o2`. This predicate is defined since we don't want to interact with an object that has objects on top of itself. If we grasp it, the object above will likely fall corrupting in this way the scene. That behaviour is undesired since a human would not grasp the bottom object without first grabbing the one on the top. Similarly for the pushing action, when an object with objects on top of itself is pushed, they could fall or collide with other objects. Vice versa if it was on top of other objects.
- **block_grasp:** (`block_grasp o1 o2`) means that object `o1` obstacles object `o2` to be grasped. Once we are sure that an object has no objects on top of it we have to check if it can be grasped, that is if the gripper will collide with some objects attempting to grasp the desired one. With this predicate the planner knows that the robot has first to interact with those objects before to grasp the desired one.
- **block_dir1, block_dir2, block_dir3, block_dir4:** (`block_dir1 o1 o2`) means that object `o1` obstacles object `o2` to be moved along direction 1. We will consider 4 possible pushing directions(defined in Chapter 4) per object. Being observant to the philosophy of human-inspired actions, we avoid collisions when we push an object, to

do that we translate the object along the pushing direction, for a certain length, and check for collision. Moreover, to push an object the end effector has to be put in the opposite side with respect the pushing direction, so an object cannot be pushed along a certain direction even in the case the gripper collides with an object. Therefore if an object cannot be moved along a certain direction because it would collide with some objects, or the gripper would collide, those have to be moved or grasped.

- **ik_unfeasible_dir1, ik_unfeasible_dir2, ik_unfeasible_dir3, ik_unfeasible_dir4, ik_unfeasible_grasp:** (*ik_unfeasible_dir1 o1*) means that the inverse kinematic to push the object *o1* along direction 1 has no solution. Similarly for the other actions. This predicates are added in order to consider the geometrical constraints regarding the working space of the robot.

Grasping Action The **preconditions** to grasp a certain object are:

- no object stands on top of it,
- no object collides with the gripper attempting to grab the desired object,
- the inverse kinematic has solution.

The **effects** of the grasping action are:

- the grasped object is dropped into the bin,
- the grasped object no more blocks other object to be pushed or grasped,
- if the grasped object was on top of other ones, it is no more on top of them.

Pushing Action The **preconditions** to push a certain object, along a certain direction, are:

- no object stands on top of it,
- the manipulated object is not on top of other objects,
- no object collides with the manipulated one, or with the gripper, when pushed,
- the inverse kinematic has solution.

In particular we defined 4 pushing action, one per pushing direction. The symbolic planner is not able to capture all the geometrical information of the problem through symbolic predicates, therefore it is not able to predict the future position of the manipulated object, and so the future states. This problem was handled by considering the object to be pushed far enough in such a way it is singulated from the other ones. Therefore the **effects** of this action are:

- the manipulated object no more blocks other objects to be pushed or grasped,
- the other objects no more block the manipulated object to be pushed or grasped.

Backtracking The geometrical constraints related to the inverse kinematic of the robot have been added in the problem description but their computation is quite expensive. Moreover, as the reader will see in Chapter 4, the actions are defined by more poses, so the inverse kinematic is even more expensive, usually it is a matter of 2 or 3 seconds per action. Computing it for each possible action (we have 5 actions in total, one grasping action and 4 pushing actions) for each object would make the computation of the predicates too expensive making the planning system quite slow. Usually the objects are inside the working space of the robot and the computation of the `ik_unfeasible` predicates is usually unnecessary. To overcome this we used the **backtracking** technique[7]. Backtracking is based on updating the initial state, used to get the current plan, considering the geometrical constraints, and plan again with the updated initial state obtaining in this way a different plan that could be feasible. The obtained plan will still be evaluated and backtracked until a feasible plan is obtained or there exists no plan. Despite this, planning symbolically is very fast therefore replanning several times is not a problem. With this method the inverse kinematic will be solved only for the action we want to execute, the first one of the plan, and no time is wasted in computing the inverse kinematic for unnecessary actions. If executing the first plan's action is not possible the equivalent `ik_unfeasible` predicate is updated. The pseudo algorithm to get a plan is shown in Algorithm 1.

Algorithm 1 Planning procedure with backtracking.

Input: initial state s_0^d and goal state G^d .

Output: a feasible plan or not plan at all.

```

procedure GETPLAN( $s_0^d, G^d$ )
repeat
    plan = GETFASTDOWNWARDPLAN( $s_0^d, G^d$ )
    if  $\neg$ ISPLANEMPTY(plan) then return NULL
    end if
    action = GETFIRSTACTION(plan)
    success = HASIKSOLUTION(action)
    if  $\neg$ success then
         $s_0^d$  = UPDATEINITIALSTATE(action)
    end if
    until success return plan
end procedure

```

It is possible that an object cannot be grasped on a certain pose, because the inverse kinematic has no solution, but it can be moved in a new pose in which it can be grasped. Therefore the pushing actions also will include the effect that the object of interest will have a solution of the inverse kinematic in the new pose. This is usually a rare case but it might happen. Of course, since the planner has no geometric information, it does not know in which pushing direction the object can be reach a position where it can be grasped, therefore this will be a totally random choice up to the planner. In the case the object is in a pose where it cannot be neither grasped or pushed because of the inverse kinematic, first the planner will return as a solution to grasp it, then it will replan and the solution will be to push it in one direction and grasp it, and so until no action can be executed and there exist no solution for the plan. In the best case, the planner will chose to push the object in direction in which the next pose will be more likely a graspable one.

A situation in which the backtracking is useful is shown in Figure 3.1. Accordingly to our strategy, the robot cannot grasp the black or red box because the gripper would collide, and the same fto push them. It has to interact with the white cup in order to

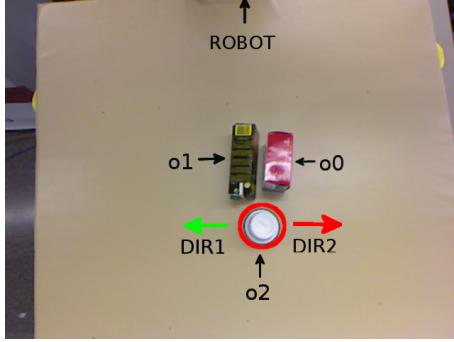


Figure 3.1: Unfeasible plan due to geometrical constraints. In this case the planner returns as action to execute grasping or pushing away the white cup (highlighted by a red circle) but it is out the configuration space of the robot and there exist no plan for that problem.

make space to move the other objects and then grasp them. The planner first returns the following plan: (`grasp o2`), (`push_dir1 o0`), (`grasp o1`), (`grasp o0`). It tries to solve the inverse kinematic for the (`grasp o2`) action, but it finds no solution. Therefore the predicate (`ik_unfeasible_grasp o2`) is added and the planner replans. The new plan is: (`push_dir1 o2`), (`grasp o2`), (`push_dir1 o0`), (`grasp o1`), (`grasp o0`). In this case it tries to get a solution for the inverse kinematic for the pushing action but it finds no solution, so the predicate (`ik_unfeasible_dir1 o2`) is added. And so on until the planner finds that all the possible actions for a feasible plan have no solutions for the inverse kinematic. In this way the planner includes the ability to understand also why no plan exists for a certain problem.

It may happen that the object outside the working space of the robot blocks the execution of the task because it is impossible to achieve the goal to grasp all the object since one is outside the working space. This happens when all the `ik_unfeasible` predicates for an object are set to true. When this situation occurs that object is removed from the goal.

PDDL For clarity purposes, the PDDL syntax of the described actions is here reported. For the *grasping* action its PDDL syntax is shown in listing 3.1.

```

1  (:action grasp
2    :parameters (?o - obj)
3    :precondition (and
4      ; grasp it if the IK has a solution
5      (not (ik_unfeasible_grasp ?o)))
6      ; grasp it if there are no objects on its top
7      (not (exists (?x - obj)(on ?x ?o)))
8      ; grasp it if there is no object that blocks it
9      ; to be grasped
10     (not (exists (?x - obj)(block_grasp ?x ?o))))
11   :effect (and
12     ; the object "o" is grasped
13     (grasped ?o)
14     ; if the object was on top of other ones now it is
15     ; no more on top of them
16     (forall (?x - obj)
17       (when (on ?o ?x)(not (on ?o ?x))))
18     ; the grasped objects no more blocks other objects
19     ; to be pushed or grasped
20     (forall (?x - obj)
21       (and
22         (when (block_grasp ?o ?x)(not (block_grasp ?o ?x)))
23         (when (block_dir1 ?o ?x)(not (block_dir1 ?o ?x)))
24         (when (block_dir2 ?o ?x)(not (block_dir2 ?o ?x))))
```

```

25      (when (block_dir3 ?o ?x)(not (block_dir3 ?o ?x)))
      (when (block_dir4 ?o ?x)(not (block_dir4 ?o ?x))))

```

Listing 3.1: PDDL syntax of the grasping action

For the *pushing* action its PDDL syntax is shown in listing 3.2.

```

1  (:action push_dir1
2    :parameters (?o - obj)
3    :precondition (and
4      ; push it if the IK has a solution
5      (not (ik_unfeasible_dir1 ?o)))
6      ; push in direction 1 only if there are no
7      ; objects that block it along that direction
8      (not (exists (?x - obj)(block_dir1 ?x ?o)))
9      ; push it if it has no objects on top of it
10     ; and if it not on top of other ones
11     (not (exists (?x - obj)(on ?x ?o)))
12     (not (exists (?x - obj)(on ?o ?x))))
13   :effect  (forall (?x - obj)
14     (and
15       ; once pushed the object is no more blocked in any direction
16       ; and it no more blocks other objects to be moved
17       (when (block_dir1 ?o ?x)(not (block_dir1 ?o ?x)))
18       (when (block_dir2 ?o ?x)(not (block_dir2 ?o ?x)))
19       (when (block_dir3 ?o ?x)(not (block_dir3 ?o ?x)))
20       (when (block_dir4 ?o ?x)(not (block_dir4 ?o ?x)))
21       (when (block_dir1 ?x ?o)(not (block_dir1 ?x ?o)))
22       (when (block_dir2 ?x ?o)(not (block_dir2 ?x ?o)))
23       (when (block_dir3 ?x ?o)(not (block_dir3 ?x ?o)))
24       (when (block_dir4 ?x ?o)(not (block_dir4 ?x ?o)))
25       ; once pushed it can be grasped and it no more
26       ; blocks other objects to be grasped
27       (when (block_grasp ?x ?o)(not (block_grasp ?x ?o)))
28       (when (block_grasp ?o ?x)(not (block_grasp ?o ?x)))
29       ; if before it cannot be grasped because of
30       ; the IK, we consider that the IK now has solution
31       (when (ik_unfeasible_grasp ?o)(not (ik_unfeasible_grasp ?o))))))
32   )

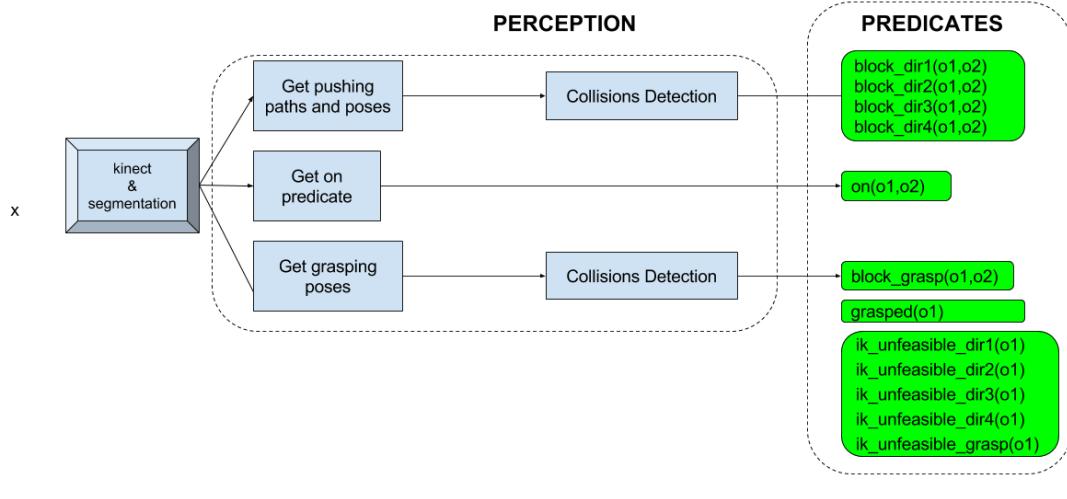
```

Listing 3.2: PDDL syntax of the pushing action along direction 1

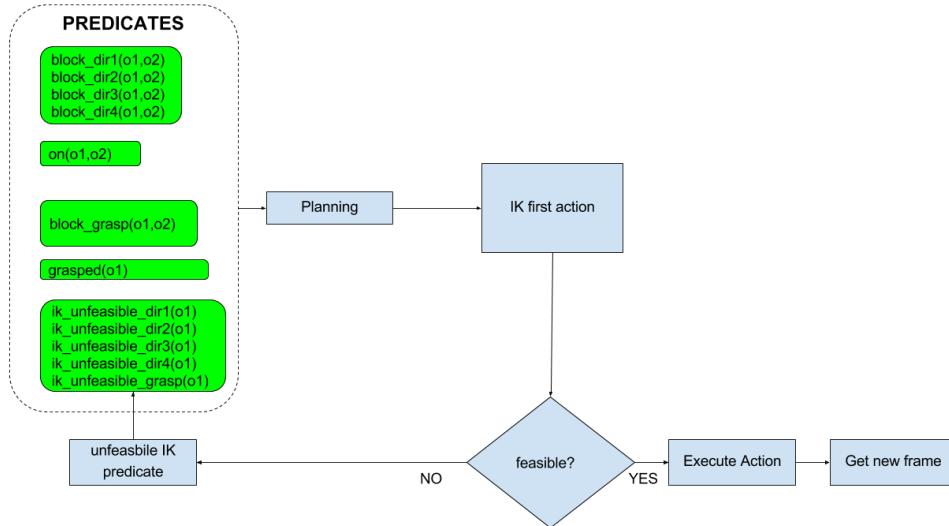
With such predicates the planner has all the basic information about the location of the objects on the scene. The symbolic predicates try to simplify the problem translating the geometric properties of the objects in symbolic ones, relaxing the planning phase. It is important to point out again that the system is deterministic, meaning that all the actions are supposed to give the resultant state with a probability of 1. Clearly the biggest uncertainty is related to the pushing action; the method used to select the pushing directions does not take into account reliably the geometry of the object and the trajectory will be unlikely the desired one but a similar one. Overall the planner is considering to push the manipulated object at infinity to isolate it, and that's false. This is another uncertainty of the pushing action due to the lack of geometrical information.

After the execution of an action the planner gets a new depth image from the Kinect, it segments the scene, it recomputes the predicates and replans. In this way the planner considers a totally new problem and all the uncertainties associated to the previous plan will be solved by the current one.

In order to provide a graphical scheme to the reader the perception and planning pipeline is depicted in Figure 3.2. It can be appreciated the two fold strategy of the algorithm, the



(a) Perception Pipeline



(b) Planning Pipeline

Figure 3.2: Perception and planning pipeline

first stage (Figure 3.2a) is devoted to get the predicates from the Kinect sensor and the second one to evaluate the feasibility of the returned plan and execute it (Figure 3.2b).

How the predicates are computed will be discussed in detail in the next chapter.

4. Implementation

In this chapter the implementation is discussed, presenting how the objects are recognized and how the symbolic predicates are obtained.

4.1. Object Localization

For the planner, to know how to move the object, is fundamental knowing where the objects are on the scenario, therefore they need to be detected. More correctly, they need to be segmented since the algorithm is dealing with unknown objects, it is not going to recognize an object as a particular one, but it segments the objects in the scene.

This stage is composed of two steps:

1. Detecting the tabletop objects
2. Segmenting the tabletop objects

The depth sensor is recording a depth map of a table, the algorithm has therefore to detect first the table, and so the objects that stands on top of it, and then segmenting them. We don't want to segment the entire image, if so, the table would be segmented as an object and the floor as well.

4.1.1. Tabletop Object Detection

The strategy for the tabletop object detection phase is composed of 3 different steps:

1. **Table plane estimation** (by RANSAC): the points of the table are detected estimating a plane in the point cloud, all the points which belong to such a plane are the points of the table.
2. **2D Convex Hull of the table**: having the points of the table a 2D convex hull is computed in order to get a 2D shape containing those points.
3. **Polygonal prism projection**: all the points are projected on the table plane previously estimated and all the points which projections belong to the 2D convex hull are considered to be points of tabletop objects.

The steps of this tabletop object detection algorithm are described in Figure 4.1 for the point cloud¹ in Figure 4.1a.

¹Point cloud taken from the Object Segmentation Database (OSD)
<http://users.acin.tuwien.ac.at/arichtsfeld/?site=4>

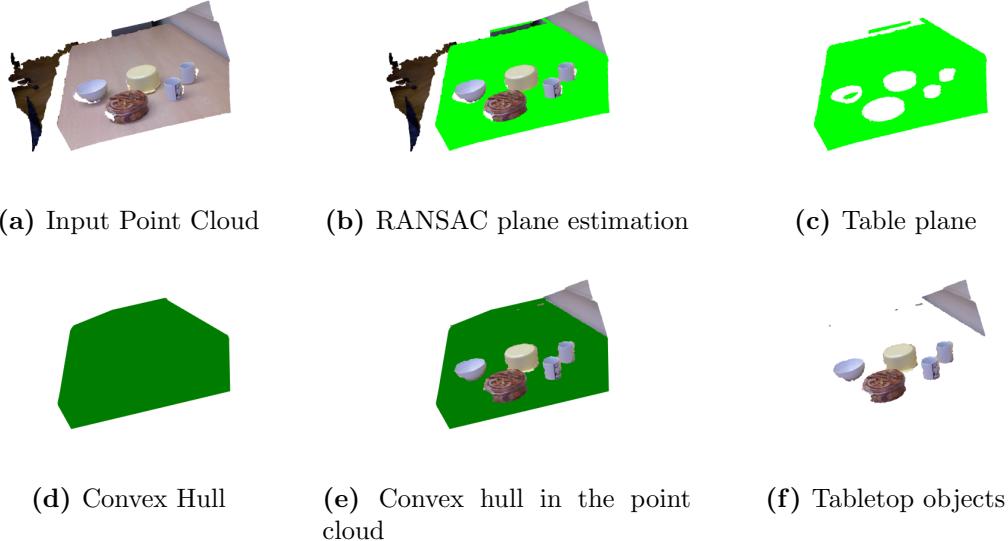


Figure 4.1: Object Segmentation: Given the point cloud (a), the estimated table’s plane is obtained (b and c), its convex hull is extracted (d and e), and the tabletop objects are obtained by a polygonal prism projection (f).

4.1.2. Object Segmentation

Once the objects on the table are detected the following phase is to segment them in order to get a point cloud per object.

Supervoxel

For their segmentation the supervoxel concept is used. A supervoxel is a group of voxels that share similar characteristics.

In this work the supervoxels are computed with the Voxel Cloud Connectiviy Segmentation (VCCS) algorithm [34], which was proposed in 2014 and gives good improvements with respect to previous state of the art methods, and still it is able to be used in online applications.

The algorithm works in 4 main steps:

- Voxelizing the point cloud
- Creating an adjacency graph for the voxel-cloud
- Creating seeds for the initial supervoxels centres
- Each seed is described by 39 features that describe spatial coordinates, colors and local surface model properties. Then a distance metric based on these features is defined.
- Clustering the voxels into supervoxles iteratively by means of the distance metric, the adjacency graph, and the search volume of the supervoxel.
- Once the search of all supervoxel adjacency graphs has been concluded, the centres of each supervoxel is updated by taking the mean of all its constituents.

Local Convex Connected Patches Segmentation

Once the supervoxels of the objects are computed, they can be clustered in order to segment the objects. Papon et al. also proposed a segmentation algorithm based on their supervoxel technique, called *Local Convex Connected Patches Segmentation* (LCCP) [36]. This algorithm permits to segment objects by clustering together adjacent convex supervoxels. The algorithm is quite simple but very good for segmentation of objects that have convex shapes, in Figure 4.2 the algorithm is briefly described.

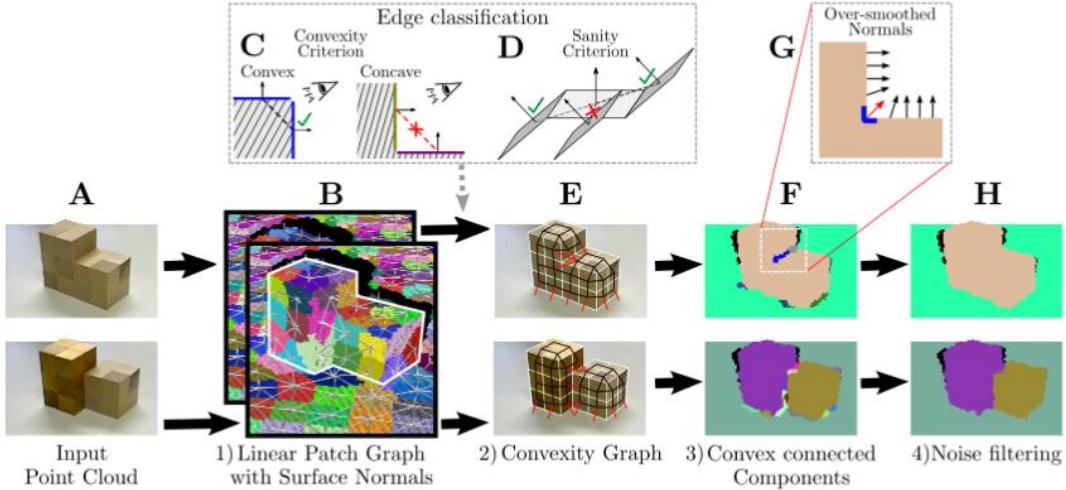


Figure 4.2: LCCP algorithm's structure. Reproduced from [36]

It clusters all the adjacent convex supervoxels (patches) using two criterion:

- Extended criterion: to consider two adjacent patches convex, both must have a connection to a patch which is convex with respect both patches
- Sanity Criterion: check if the adjacent patches which can be considered as convex present geometric discontinuities (see point D of Figure 4.2), in this case they are not considered as valid to form a cluster.

Then, due to the smoothed normals that could appear in some edges of the objects (point G Figure 4.2), the algorithm merges the clusters that are composed of few supervoxels to the biggest adjacent cluster.

By tuning properly the parameters of the segmentation algorithm the objects can be correctly segmented obtaining for one of them a point cloud. Two examples of the segmentation algorithm for a cluttered scene are depicted in Figure 4.3.



Figure 4.3: Example of segmentation results.

Note that the algorithm works mainly considering geometric properties, and not the color of the pixels. Considering the colors could lead to worst segmentation results for our case of studio since many objects have no only one color.

A color based segmentation could segment a draw on an object, or a small part of the object, as a different object, but this, accordingly to the strategy we are going to use (see next sections), would lead to an unfeasible problem. For instance, in Figure 4.4 is shown a box with a green stripe on its top surface, a segmentation algorithm based also on colors could lead to segment the green stripe as another object, and the result is that it is impossible to grasp the green stripe without collide with the box. Vice versa, any interaction with the box is impossible without having a collision with the green stripe. This is the main reason the segmentation we used is based only on geometric features (The LCCP algorithm has the option to include color importance but we neglected it).

Talk about filtering

Figure 4.4: Box with a green stripe.



4.2. Background

In this section will be presented some concepts that will be used to execute the actions and to compute the predicates.

Principal Components The first concept to know is the concept of principal directions. The principal direction of an object is its principal axis which is defined as any of three mutually perpendicular axes about which the moment of inertia of a body is maximum. For instance, for a rectangular object its principal direction is the axis aligned with its longest dimension.

To obtain the principal axis the principal component analysis (PCA) [20] technique is used. This technique is a common statistical procedure that uses orthogonal transformation to convert a set of observations of possibly correlated variables into a set of values of linearly uncorrelated variables, which are called principal components. The transformation is defined in a manner that the first component has the largest variance, the second has the second largest variance and so on. The principal components are orthogonal because they are the eigenvectors of the covariance matrix, which is symmetric. An example of the principal components for a 2D data set is depicted in Figure 4.5a². The principal components are computed through the covariance matrix of the set of observation, and its eigenvectors $\bar{\lambda}_v$ represent the principal components while its eigenvalues λ represent the variance of the data set along the principal component $\bar{\lambda}_v$.

A generic point cloud can be seen as a set of observations and the PCA can be directly applied in to the object's point cloud to retrieve its principal components, or how called in this work principal directions. In Figure 4.5b the principal directions of a generic object are illustrated.

Projection onto a plane We will see later that several things need the concept of the projections of a point into a plane. Considering a point $p = x_p, y_p, z_p$ and a plane P defined by the following equation

$$ax + by + cz + d = 0 \quad (4.1)$$

²Image taken from https://en.wikipedia.org/wiki/Principal_component_analysis

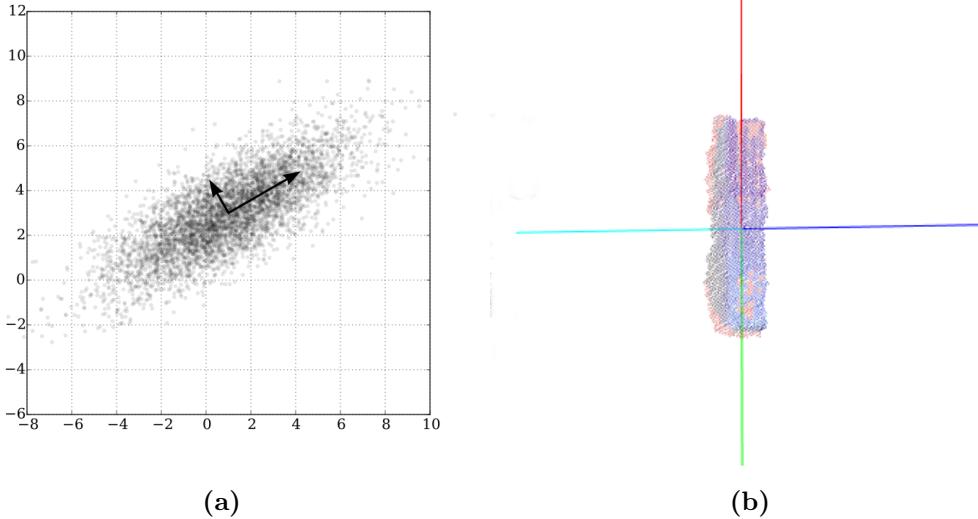


Figure 4.5: Principal Components Analysis - In Figure 4.5a PCA for a standard 2D set of observations. In Figure 4.5b results of the PCA for a rectangular segmented object. The green, red lines refers to different ways of the first principal components, while blue and cyan lines refers to different ways of the second principal directions. The third one would be orthogonal to the first two principal components.

the projection p_P of point p onto the plane P is given by the following set of operations:

- Given the origin point $P_0 = (x_0, y_0, z_0)$ of the plane, which can be calculated by arbitrary x_0 and y_0 coordinates as

$$z_0 = \frac{-1}{c}(ax_0 + by_0 + d),$$

- calculate the coordinates of P_0 with respect the point p

$$P^p = p - P_0,$$

- then calculate its projection λ_p onto the plane normal $n = (a, b, c)$

$$\lambda_p = \bar{n} \cdot p^{plane},$$

- compute the coordinates of the projection onto the plane of point p with respect to the point itself:

$$p_P^p = -\lambda_p \bar{n}$$

- translate it by the coordinate of the points in order to have the projection of the point expressed with respect the original reference frame:

$$p_P = p_P^p + p.$$

Rotation Matrices Rotation matrices are matrices which express a rotation between two reference frames. Given two frames $\{A\}$ and $\{B\}$, and the rotation matrix ${}^A_B R$ that defines the rotation of $\{B\}$ relative to $\{A\}$ then a point ${}^A P$ with respect frame $\{A\}$ is given by ${}^A P = {}^A_B R {}^B P$, where ${}^B P$ is the same point relative frame $\{B\}$.

Having a desired frame $\{B\}$ defined by axis ${}^A\hat{X}_B$, ${}^A\hat{Y}_B$ and ${}^A\hat{Z}_B$, where ${}^A\hat{Y}_B$ is the y axis of frame $\{B\}$ relative to frame $\{A\}$, the rotation matrix between $\{B\}$ and $\{A\}$ is defined as

$$R_B^A = \begin{bmatrix} {}^A\hat{X}_B \\ {}^A\hat{Y}_B \\ {}^A\hat{Z}_B \end{bmatrix}$$

To transform any object, such as the gripper mesh model, to frame $\{B\}$ then the following homogeneous transform is applied:

$$H = \begin{bmatrix} R_B^A & {}^A B_O \\ \vec{0} & 1 \end{bmatrix}$$

where $R_A^B = R_B^A{}^\top$ and ${}^A B_O$ is the origin of frame $\{B\}$ relative to $\{A\}$. In this way, having some axis that defined our new reference frame, we can transform the gripper model in such a way its closing point is in the origin of the new frame and its orientation is the one defined by the new reference frame.

Axis Aligned Bounding Box (AABB) Axis aligned bounding box is a bounding volume defined with respect the reference frame of an object. It is therefore aligned to the objects principal components. A bounding volume is a closed volume that completely contains the object³. After a transformation of the object from its frame and the world frame the dimension of the bounding box are obtained by computing the maximum and minimum coordinates of the point cloud. In this way it is possible to have an approximation of the length, width and height of an object.

Convex Hull A convex hull of a point cloud P is the smallest 3D convex set that contains P . In Figure 4.6 an example of the convex hull for a point cloud is shown. The vertices are

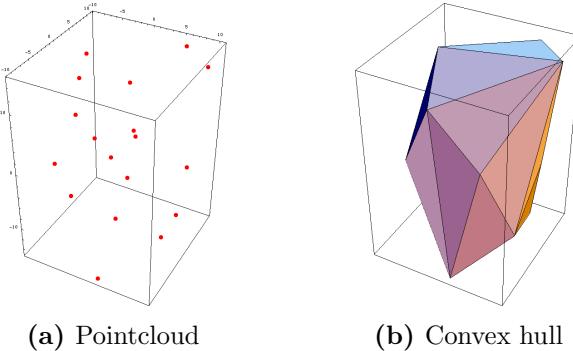


Figure 4.6: Convex hull example

first detected and then connected among them by means of triangles. In this way a triangle mesh is associated to the convex hull.

Collision Detection To understand if an object impedes a certain action, such as the pushing along a direction, we have to check if along the desired trajectory the pushed object will collide with a certain one. The collision detection is therefore a crucial step for the predicates computation. There exist different techniques to assert if two objects are

³https://en.wikipedia.org/wiki/Bounding_volume

colliding, all of them need a representation of the object, which could be for example a basic shape or a more complex as an octree.

The mesh shape has been thought to use since it can be directly obtained from a *convex hull*.

Given two objects A and B and their configurations \mathbf{q}_A and \mathbf{q}_B , the discrete collision query returns a boolean value about whether two objects collide or not. Two objects collide if

$$A(\mathbf{q}_A) \cap B(\mathbf{q}_B) \neq \emptyset$$

The collision detection will be used to understand if in a given pose \mathbf{q} the object A will collide with the other objects in the scene.

To make the collision detection most collision libraries, before to use complex algorithm to detect collision between two shapes, they first check if the AABB of the objects intersect, if they don't the objects surely don't collide. If their AABB intersect the objects might collide.

Objects Modelling The triangle mesh needed for the collision checking can be directly retrieved by the convex hull. But there is a consideration to take into account. The Kinect, given its pose, is only able to see mainly the top surface of the objects and not all the sides, and therefore we cannot apply directly the convex hull algorithm to the detected surfaces. If we would apply the convex hull on an object's surfaces, we will have likely the situation in which a surface S transformed in a configuration such that it will be above or below another one, and the collision detection will detect no collision since the two surfaces are not colliding. This is because we are totally missing the rest of information about the geometry of the object since we only know the objects surfaces.

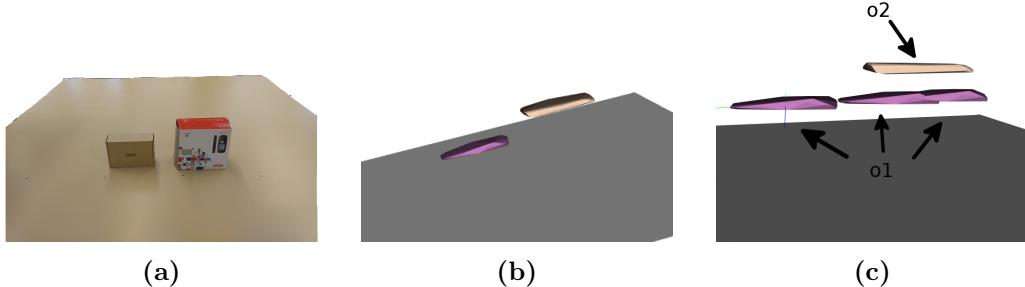
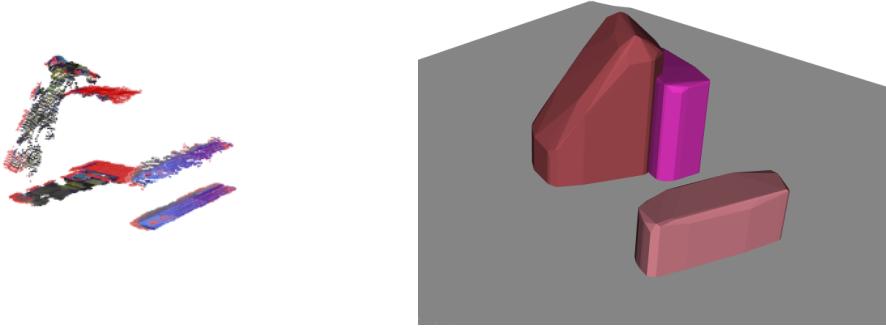


Figure 4.7: Convex hulls and collision detection using the segmented objects retrieved by the LCCP segmentation algorithm. The gray surface represents the plane's 2D convex hull. In Figure 4.7b it is possible appreciating that we miss the information about the hidden part of the object. In Figure 4.7c a collision detection example is depicted. The convex hull is translated along a direction and no collision is detected with any possible translation since the two convex hull will not intersect.

From the Kinect's point cloud also the table plane is known, so all the information we have are: the table plane model and the segmented objects (mainly the top surfaces). If a human would be in the same pose of the Kinect, looking at the table, it will imagine that the objects are not floating surfaces, and he/she will deduce the objects shape from the shape of the top surface. Since in this work we assumed to work with parallelepipeds, the sides of the objects can be easily deduced by projecting the top surface's edges to the plane and then filling the missing object's sides with points. To do that we have to detect the top surface's edges. A more trivial method is directly projecting all the points of the surfaces onto the table plane and then apply the convex hull algorithm to the resulting point cloud given by the sum of the top surface and its projection. In this way the missing sides are indirectly retrieved by the convex hull.



(a) Surfaces projection

(b) Resulting convex hull

Figure 4.8: Convex hull of the objects using their projections onto the table plane.

4.3. Actions Execution

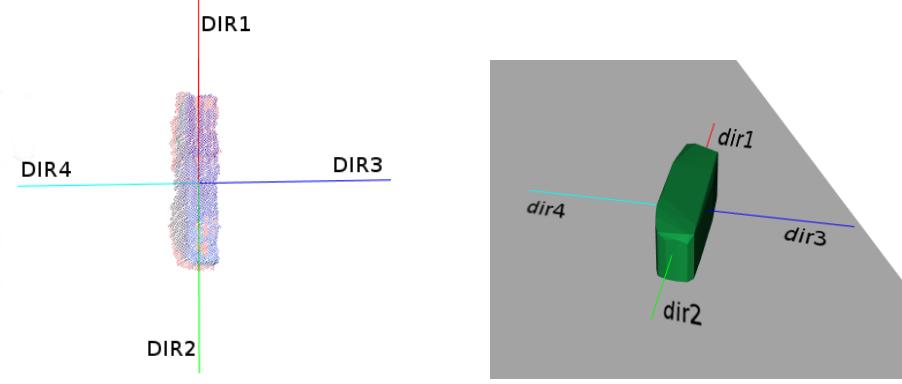
In this section the way the actions are executed is discussed, since the predicates computation depends on the way we decided to execute the actions.

4.3.1. Pushing

Pushing is a difficult action to execute when the goal is to move one object from one pose to another one. At the best of my knowledge there are still no methods in state of the art that faced this problem. The majority of pushing actions in object manipulation have the aim to interact with the objects in order to move them and validate the segmentation [23] [21] [22], without taking care about the final position of the objects. Hermans er al. [19] presented a novel algorithm for object singulation through pushing actions, here the pushing actions have the aim to separate objects in cluttered scenes but they are not interested in singulate tidily the objects but just to find a feasible interaction to singulate them. This means that the pushing action could move several objects and the pushing direction is not chosen accordingly the object to push but accordingly to the plane that separate the objects. This results in a pushing action without any guarantee that the pushed objects follows the desired path. Grasping, although is a difficult task, has already an advanced state of the art regarding finding grasping pose also for complex objects. Somehow grasping is easier because it does not consider the whole object geometry but it is focused on looking for local futures that can be interesting for grasping the object (i.e. antipodal points). When we push an object all the object geometry has to be taken into account in order to find a suitable pushing direction. Moreover, humans push objects with a certain initial direction chosen accordingly the object geometry and then adapt the pushing direction accordingly to the error with respect the desired one. This would imply the implementation of a controller which is out the scope of this thesis.

We are considering to work with objects with simple shapes, such as parallelepipeds. A human would push such object mainly accordingly its principal axis and the one orthogonal to its (i.e. the first 2 principal components of Figure 4.5b), but he/she also could push it along the diagonal thanks to the several amount of freedom of the hands. Inspired by this consideration, we decided to consider as possible directions to push an object its first two principal components, the third is orthogonal to the object and it is not of interest for pushing. In particular, there are two senses for each direction, so in total we have 4 possible pushing directions per object, as depicted in Figure 4.9.

Another things to take into account is that the principal directions are not parallel to



(a) Pushing directions computed using the segmented surface seen by the Kinect.
(b) Pushing directions associated to the object's convex hull.

Figure 4.9: Example of pushing directions.

the table plane. The Kinect could see some sides of the objects, and these sides will affect the direction of the principal components. An object which stands on top of a table will be obviously pushed along a direction parallel to the table. As previously described in Section 4.2 a point can be projected on to the table plane, but a point and a vector (the principal components are vectors) have the same representation, and therefore we can apply exactly the same equations to obtain the projections of the principal components onto the table plane. So the pushing directions considered are not the one obtained by the PCA but their projections.

Next, having some pushing directions along with the robot will push the objects, the pose of the end effector (i.e. the gripper) to push the object has to be decided. The pose has been chosen accordingly to the shape of the end effector.

In Figure 4.10 is possible observing the profile of the gripper mounted to the last link of the robot, highlighted by the blue color. Such last link has a circular shape and the gripper's deep is less than the last link. It is undesirable pushing an object with a circular, or spherical, shape for the end effector because there is more probability to make the object no follow the desired path. The gripper has no a circular shape and it is all symmetric, this make it suitable to push an object with a certain stability, i.e. make the object follow the desired path, during the action. As final consideration, we don't want that the last link of the robot touch the objective object since we want a pushing action accurate as most as possible.

Knowing also the height of the objects retrieved by its AABB, it is possible having a pose for the gripper is such a way that the end effector does not touch the object. The gripper's pose, relative the object, is computed in manner to locate the red point of Figure 4.10 to be at the same height of the object. In this way the fingers will fully touch the object during the pushing action. Moreover, to make easy for the robot reaching the pushing pose, it was defined to be a certain distance from the object (in our experiment it was set to 5cm). It would be difficult to reach a pose which is tangent to an object without colliding with it.

Due to the limited opening width of the gripper (8 centimetres) the object the robot is going to manipulate have small width. This means that when pushing along the principal



Figure 4.10: Gripper's profile view and robot's end effector.



Figure 4.11: Profile view of a desired pose for pushing an object.

axis, the object's width is likely small. Such a situation is depicted in Figure 4.12a. Pushing in such a way the gripper will likely push also the black juice box. Therefore when pushing along the principal axis the pose is chosen to be the one in Figure 4.12b. Of course is more stable a pushing pose like the one in Figure 4.12a, since the contacts point (the fingers) are more distant. For this reason that pose is used only when pushing the objects along a direction orthogonal to the principal axis.



Figure 4.12: Possible pushing poses for push an object along its principal axis - In Figure 4.12a the closing direction of the gripper is orthogonal the pushing direction, and for the case depicted in the figure the gripper will likely push also the black juice box. In Figure 4.12b the closing direction of the gripper is parallel to the pushing direction.

Having the projections of the principal components, the table normal and the gripper closing point coordinates it is possible defining a transform defined by the following rotation matrix and translation vector:

$$R_{dir1} = \begin{bmatrix} dir2_X & dir2_Y & dir2_Z \\ dir4_X & dir4_Y & dir4_Z \\ n_x & n_y & n_z \end{bmatrix}^\top \quad T = \begin{bmatrix} c_x \\ c_y \\ c_z \end{bmatrix} \quad (4.2)$$

where $dir1_X$ refers to the x coordinate of the vector that defines the direction 1, n is the table normal and c is the desired tool central point. Eq. 4.2 refers for a pushing along direction 1, where the gripper's closing direction has to be parallel to the pushing direction. The rotation matrix R_{dir3} for pushing along direction 3 would be the same of R_{dir1} .

As previously said, the planner consider to push the objects at infinity, in reality the robot has to push the objects for a finite length. Since the planner has no geometric information, the pushing length is chosen accordingly the dimension of the AABB relative

to the pushing direction. For instance, if the robot is going to push the object 1 along direction 1 the length l of the pushing action is $l = k \cdot \text{AABB}_1.length$, where k is a gain factor (1 in our experiments) called *pushing step*. This is a big limitation since the robot will push an object accordingly to the manipulated objects, not to surrounding ones.

The first pose is the pushing pose while the last pose is the pushing pose translated along the pushing direction by the length of the push. To retrieve the path we consider the total length and we discretize it by n points having in this way $n + 2$ poses (2 because of the pushing pose and the final pose). For each pose the inverse kinematic is done. In this way we obtain a discrete path.

When the robot approaches the pushing pose it could be that it collides with other objects. It would be suitable to use *MoveIt!* which can built an octomap representation of the scene and find a path avoiding collision with the scene. Although this technique is very powerful and correct for this application is quite time consuming. We avoid this by simply considering a pre pushing pose which has the same pose of the pushing pose but translated, accordingly the table normal, 10 centimetres from the pushing pose. After the execution of the pushing action the robot goes to its *home* pose (depicted in Figure 1.5a) in order not to fall inside the Kinect's view. When it goes to home it might happen that it collides with some objects, so also for the final pose with consider another one translated, accordingly the table normal, 10 centimetres from the last pose. In this way the pushing trajectory is defined by a total of $n + 4$ poses.

4.3.2. Grasping

There exist an advanced state of the art regarding grasping. Despite this all the techniques of grasping are very computationally expensive. Many of them rely on the identification of the shape of the objects and then a set of pre-built grasping poses is returned and they can be very robust as the method proposed by Brook et al. [8]. Other techniques rely on the identification of local features which can state if a grasping pose is feasible or not. Two very goods grasping planning algorithms, which deal with novel objects, published past year are AGILE [37] and HAF [17], despite this how shown in [11] they are not so robust and they are computationally expensive and not suitable for this thesis. In order to have a faster planning algorithm we considered a very simple approach to grasp the objects, which is suitable only with the kind objects we are going to interact with. Despite this the planner present by this thesis can be directly integrated (with just few modifications) with several grasping algorithms.

The idea is to grasp the object in manner that the gripper's closing direction is orthogonal the principal axis of the object. The approaching direction of the gripper is given by the third principal component of the object. Then the gripper also is centred accordingly the centroid of the object. In this manner a single grasping pose is obtained for each object.

Add exactly how the gripper position is computed

To grasp the object also the robot needs a pre grasping pose, if not the gripper will collide with object attempting to reach the grasping pose, moving it away, and the grasp will fail. The pre grasping pose is simply defined by the grasping pose translated along its approaching vector by 10 centimetres.

The grasping action, in its total is composed by the following set of actions:

1. Reach the pre grasping pose
2. Open gripper
3. Reach grasping pose

4. Close gripper
5. Reach pre grasping pose again: this is done in order to avoid collisions between the grasped object and the other ones
6. Go to the dropping pose: the object will be dropped into a bin

4.4. Predicates Computations

In Chapter 3 the predicates used are described, in this section their computation is presented in detail.

Once the objects have been segmented, we have one point cloud per object, next we have to retrieve the symbolic predicates from the objects configuration.

4.4.1. Predicate: block_grasp

The `block_grasp o1 o2` predicate refers the fact that object `o1` impedes `o2` to be grasped. The computation of this predicate is straightforward: the mesh model of the opened gripper is transformed to the grasping pose of object `o2`, and check if it collides with the other objects. In figure 4.13 such procedure is shown and in Algorithm 2 the pseudo algorithm is described in detail.

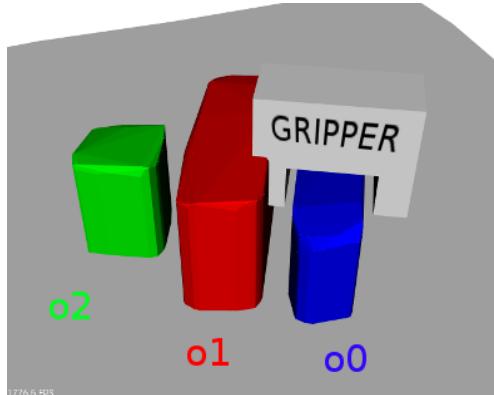


Figure 4.13: Visualization of the computation of `block_grasp` predicate for object `o0`. The opened gripped model is transformed to the grasping pose for object `o0` and it is tested if the gripper mesh model collides with the other objects, in this case it collides with `o1`.

Notice that this method implies to check for collision between the gripper and objects that might be very far from the interested object, i.e. there is no need to compute the collision detection. Despite this, as explained in Section 4.2 the most collision detection algorithm first check if the AABB of the objects intersect. This is a computationally cheap operation, and only if their AABB intersect the computationally expensive algorithm are used to check for collision. This makes the Algorithm 2 efficient and computationally not expensive. It has been observed that, in average, to compute this predicate the time is about 8 milliseconds. This is for two main reasons, the objects are far between each other and their AABB do not intersect, or in case they intersect the gripper mesh model is relatively simple compared to the convex hull of an object, this makes the collision detection algorithm faster because it has to check for the intersection of less triangles.

4.4.2. Predicate: on

The `(on o0 o1)` predicate means that object `o0` is on top of object `o1`. With the convex hull of the objects is easy to understand if two objects are one on top of the other or by

Algorithm 2 Computation of `block_grasp` predicate. O is the set of objects (convex hull retrieved with the projection onto the table plane) and G_{poses} is the set of grasping poses of all the objects.

```

function COMPUTEBLOCKGRASPPREDICATES( $O, G_{poses}$ ) block_grasp_predicates =  

    NULL  

    for all  $A = objects \in O$  do  

         $openGripperMesh \leftarrow$  TRASNFORMGRIPPERMODEL( $G_{poses}(A)$ )  

        for all  $B = objects \in O$  do  

            if  $A \neq B$  then  

                 $collision \leftarrow$  IS THERE COLLISION( $openGripperMesh, B$ )  

                if collision then  

                    block_grasp_predicates = ADDPREDICATE(block_grasp B A)  

                end if  

            end if  

        end for  

    end for  

    return block_grasp_predicates  

end function

```

checking for collision, but in this way we do not know who is above and who is below. To do this their surface projections onto the table plane are used. The research group of Artificial Intelligence and Robotics Laboratory of Istanbul Technical University, published some interesting researches suitable to the aim of this thesis. In [15] [32] [16] the authors proposed some approaches to enhance 3D recognition and segmentation results to create and maintain a consistent world model involving attributes of the objects and spatial relations among them. Their researched focused on modelling the world for manipulation planning tasks. They do not consider scene like the one of this thesis but simpler ones such as a pile of cubes above each other. What can be directly used from their work is the computation of the `on` predicate. The `on` relation for a pair of objects is determined by checking whether their projections on the table plane overlap. This predicate was not a relevant part of their work and they did not provide to much information about its computation. Therefore our implementation for the `on` predicate is based on their idea with some modifications.

Our idea is based on the fact that an object which stands on top of another one occludes some parts of the object below. While the one below does no occlude any part of the top object. Let's consider the scene in Figure 4.14a, the object o_0 occludes some parts of object o_1 . The projections P_0 and P_1 onto the table plane of o_0 and o_1 are respectively the red and green ones in Figure 4.14b. This means that the convex hull C_{P_1} of the projection o_1 intersect with the projection P_0 of o_0 , while the projection P_1 of o_1 does not interesct with the convex hull C_{P_0} of the projection of o_0 (Figures 4.14c and 4.14d).

Although this method works fine to compute the `on` predicate it has the limitation that its scope is only for objects for a rectangular shaper. An object that has a form of L has a convex hull that includes big free areas where an object could lie.

It is important to take into account also that actually the edges of the occluded parts of the below object, once projected, could be at the same position, of some projected edges of the top object. This could be dangerous for the computation of this predicate. Therefore a threshold is added. Focusing the attention on Figure 4.14c it can be appreciated that the intersection $C_{P_0} \cap P_1$ has several points, while in case the two edges relative to the occluded and occluding part have similar coordinates, the intersection $C_{P_1} \cap P_0$ would have just few points. Therefore the (`on o0 o1`) predicate is update accordingly to the following

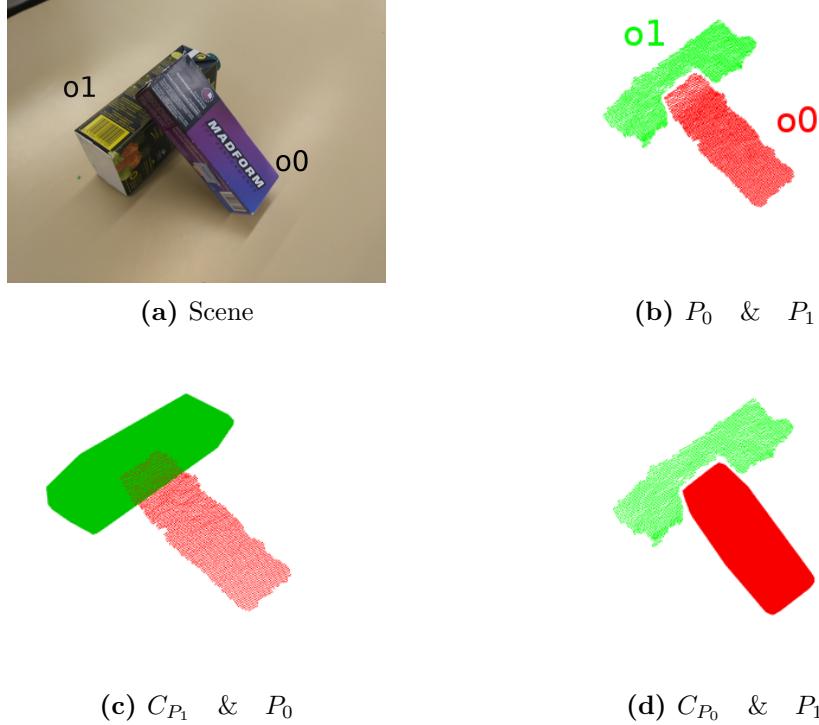


Figure 4.14: Visualization of the computation of the `on` predicate. Figure 4.14b show the real image, Figure 4.14b shows the projections of the objects onto the table plan while Figures 4.14c and 4.14d represent the two step strategy to compute the `on` predicate.

mathematical formula:

$$(\text{on } o_0 \ o_1) = \begin{cases} \text{True}, & \text{length}(C_{P_0} \cap P_1) > th_0 \wedge \text{length}(C_{P_1} \cap P_0) < th_1 \\ \text{False}, & \text{otherwise} \end{cases} \quad (4.3)$$

where $\text{length}(A)$ means the number of elements in the set A . The values of the thresholds th_0 and th_1 are determined empirically and they are $th_0 = th_1 = 100$.

The formula 4.3 is then evaluated for every possible combinations of objects, that is $n(n - 1)$ times, where n is the number of objects. Despite this, its computation is very fast. For the example in Figure 4.14 its was evaluated just 2 times since there are only 2 objects and it took 3 milliseconds, that is $\approx 1.5 \frac{\text{ms}}{\text{pair of object}}$. For instance, for a complex scene with 10 objects the total time devoted to compute this predicate would be $10 \cdot 9 \cdot 1.5 \approx 135\text{ms}$ (the \approx symbol is due to the fact that the computation time depends on the number of points of the objects and the complexity of their convex hulls).

4.4.3. Predicate: `block_diri`

The `(block_diri o1 o0)` predicate, if true, means that object o_1 impedes object o_0 to be moved along its i -th direction.

Object o_1 can impedes object o_0 to be moved along a certain direction if a collision will appear between the two objects. In order to do that, having a certain pushing length l_i for the i -th direction of object o_0 , its convex hull C_0 is translated along the considered direction until reach its final position which is $p_f = p_i + l \cdot dir_i$, where p_f and p_i are respectively the centroid at the final and initial pose. Object o_0 is going to do a path from its initial and final pose so the collision should be checked along its path. To do that we could use a

continuous collision detection algorithm or a discrete one. The continuous method suppose that given an object, its initial, its final pose and the other objects in the environment it will detect a collision if a collision occurs from the initial to the final pose. This algorithm has the advantage to be accurate but we should repeat it separately for each object since we are not interested to see if there is a collision, but we are interested in knowing what is the object it is going to collide with. We decided therefore to use the discrete strategy, that is we consider several poses between the initial one and the final one, including the final one, and for each check if it collides with the other objects.

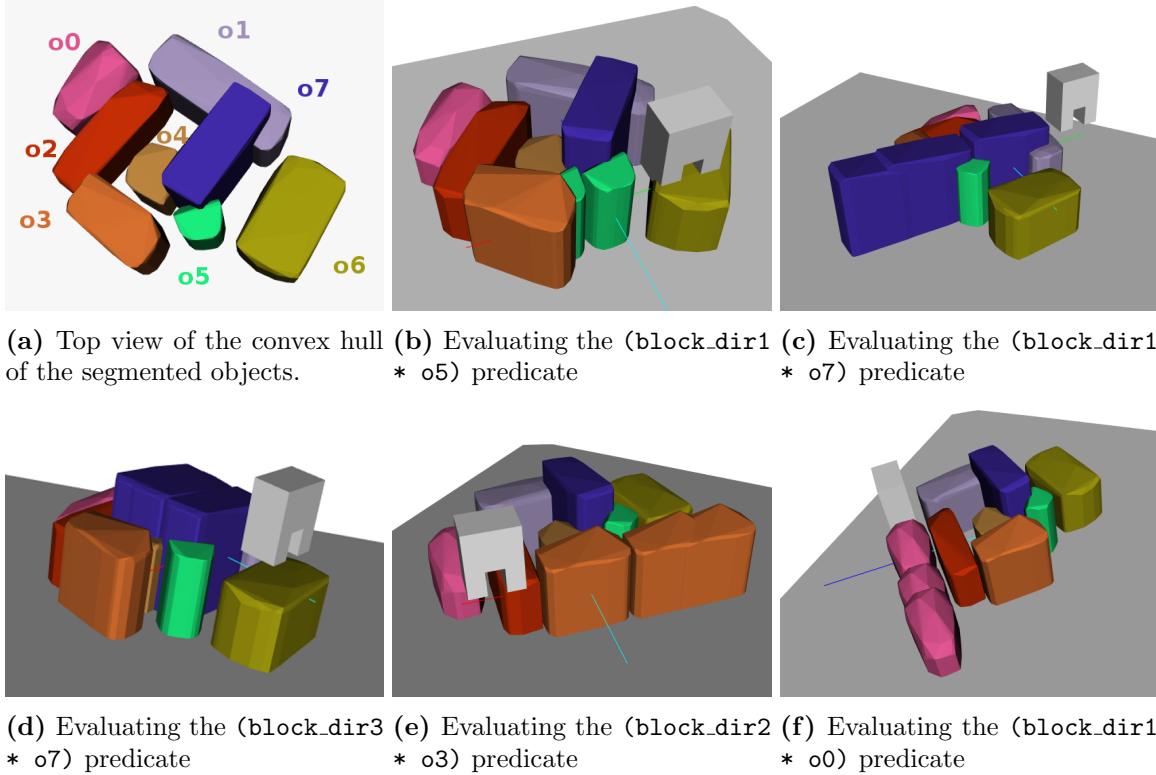


Figure 4.15: Visualization of the computation of block_dir_i predicates. "Evaluating the $(\text{block_dir}_1 * o_0)$ " means that the algorithm is evaluating for all the objects, except o_0 , if they collide with o_0 when pushed along direction 1.

Considering the kind of objects we are going to interact with, the discrete path is computed by translating the object along the pushing direction by a length equal to the AABB dimension associated to that direction, until reaching the final pose. Note that no collision detection is done for the object at its initial pose, therefore the total number of poses considered for the pushing path are $n_{\text{poses}} = \left\lceil \frac{k \cdot \text{AABB}_{\text{dimension}}}{\text{AABB}_{\text{dimension}}} \right\rceil = [k]$, where k is the pushing step defined in Section 4.3.1.

To push an object along a certain direction the robot needs to put its end effector at the opposite side of the object. Therefore object o_1 can impedes object o_0 to be moved along a certain direction also in the case the end effector cannot be put in the pushing pose because it would collide with o_1 . This computation is simply done by transforming the closed gripper mesh model to the pushing pose relative the desired pushing direction of the object to push and check for collision with the other objects, as similarly done for the computation of the block_grasp predicate.

In Figure 4.15 is shown graphically the procedure to compute the predicate.

Note that also the gripper during the pushing action will move, so ideally the collision checking should be done exactly how done for the object. We decided to neglect this and

check only for the initial pose in order to relax the planner and not to make it too much conservative. This means that during the pushing action the robot might actually move more than one object. Despite this relaxing strategy the algorithm showed to work perfectly and cases in which the gripper moved more than one object were really rare, at least with the objects we used in our experiments.

Algorithm 3 Computation of `block_dir` predicates. O is the set of objects (convex hull retrieved with the projection onto the table plane), P_d is the set of the pushing directions of all the objects, P_{poses} is the set of the initial pushing pose of all the object and P_{path} is the set of all the pushing poses relative to each direction and each object.

```

function COMPUTEBLOCKDIRPREDICATES( $O, P_d, P_{poses}, P_{path}$ )
    block_dir_predicates  $\leftarrow$  NULL
    for all  $A = \text{object} \in O$  do
        for all  $d = \text{directions} \in P_d(A)$  do
            for all  $p = \text{poses} \in P_{path}(A, d)$  do
                for all  $B = \text{object} \in O$  do
                     $A_T \leftarrow \text{TRANSFORMOBJECT}(A, p)$ 
                    if  $A \neq B$  then
                         $collision \leftarrow \text{IS THERE COLLISION}(A_T, B)$ 
                        if  $collision$  then
                            block_dir_predicates  $\leftarrow \text{ADD PREDICATE}(\text{block\_dir } B \ A)$ 
                        end if
                    end if
                end for
            end for
             $closedGripperMesh \leftarrow \text{TRASNFORMCLOSEDGRIPPERMODEL}(P_{poses}(A, d))$ 
            for all  $B = \text{object} \in O$  do
                if  $A \neq B$  then
                     $collision \leftarrow \text{IS THERE COLLISION}(closedGripperMesh, B)$ 
                    if  $collision$  then
                        block_dir_predicates  $\leftarrow \text{ADD PREDICATE}(\text{block\_dir } B \ A)$ 
                    end if
                end if
            end for
        end for
    end for
    return block_dir_predicates
end function

```

The computation of this predicate can be appreciated in detail in Algorithm 3. From that pseudo code can be noted that the complexity of this function to compute the predicate approximatively is $\mathcal{O}(n \cdot n_{\text{directions}} \cdot n_{\text{poses}} \cdot (n - 1) + (n - 1) \cdot n_{\text{directions}})$.

The computation of this algorithm is the most computational expensive of all the planner since it involves many collision detection phases. In fact the time required to compute this predicate for the example in Figure 4.15 (considering to push the objects 1.5 times the AABB dimension relative to the pushing direction) was $\approx 1.115\text{seconds}$, were $\approx 0.869\text{s}$ were dedicated to check if the objects will collide when move, while $\approx 0.239\text{s}$ were dedicated if the gripper collides with some objects.

5. Software design

Talk about ROS and real robot.

For the collision checking the Flexible Collision Library (FCL) [33] has been used. This library allows to define the collision problem in a simpler manner than other more famous collision libraries such as *Bullet* [4], and it can work with different objects shapes such as box, spheres, cone, convex, mesh and octree. The main library used in this work is the Point Cloud Library (PCL) [2], which allows some methods to create an object shape from a point cloud.

1. Nodes
2. Ros graph
3. simulation
4. PCL - FCL

6. Experiments

In this chapter the performance of the proposed task planner for the considered task is evaluated by performing several experiments with the aim to study the quality, its benefits and its limitations.

To get the quality of the task planner the strategy defined in the article [6] by Amignoni et al. In that article they proposed a general method to evaluate the quality of a task planner. First they define a functionality benchmark¹ thought 4 elements:

- Description: a high level, general, description of the functionality
- Input/Output: the information available to the module implementing the functionality when executed, and the expected outcome.
- Benchmarking data: the data needed to perform the evaluation of the performance of the functional module.
- Metrics: algorithms to process benchmarking data in an objective way.

Regarding the benchmarking of a task planner those elements are defined as:

Description: Given a initial state and a goal state find a set of actions to reach the goal.

Input/Output: The inputs are the PDDL domain and problem descriptions. The output is a plan expressed as sequence of actions.

Benchmarking data: Some data has to be supplied by the robot, including: the input given to the planner, the output (i.e., the plan) delivered by the planner, and the performance data for the planner (e.g., memory and time needed). According to the planning approach, other data might be required. Other data will be possibly collected when a plan is actually executed, including: time needed to execute each action and time needed to execute the whole plan.

Metrics: Metrics are mainly related to good plan quality and to solving time. The target of benchmarking task planning is twofold: to asses the quality of the plan produced (e.g., in terms of the cost for executing it and its likelihood to succeed) and to assess the quality of the planning process (i.e., how fast is the plan generated and how much resources are needed to do so). Possible metrics are:

- *Levenshtein distance*² is a metric that specifies the distance between two strings, in this case a plan can be though as a sequence of character, where each action

¹A functionality benchmark is a benchmark of specific robotic task. In the case of this thesis a benchmark of the task planner.

²https://en.wikipedia.org/wiki/Levenshtein_distance

with a determined object of interest is a character of the string. The Levenshtein distance between the two lists is calculated to measure the quality of the plan. Levenshtein distance $L_{A,B}(|A|, |B|)$ between two lists A and B is calculated as

$$\begin{aligned} L_{A,B}(i, j) &= \max(i, j) \quad if \quad \min(i, j) = 0 \\ L_{A,B}(i, j) &= \min(L_{A,B}(i - 1, j) + 1, \\ &\quad L_{A,B}(i, j - 1) + 1, \\ &\quad L_{A,B}(i - 1, j - 1) + 1_{a_i \neq b_i}) \quad otherwise \end{aligned}$$

where $|A|$ is the length of plan A, and $1_{a_i \neq b_i}$ is equal to 0 if the action a_i is the same of action b_i . This metric will be used to evaluate the distance between the optimal plan returned by the planner in the first frame and the real executed plan.

- Number of actions correctly performed in the correct order (e.g., number of objects delivered to their correct destination).
- Time required to construct a plan and time to perform it.

We will also evaluate the difference of executing the initial plan and executing the plan with the replanning technique in terms of time and number of correct actions. As well also the number of bad, or dangerous, object manipulations due to the lack of replanning, lack of probability and geometric constraints in the planner .

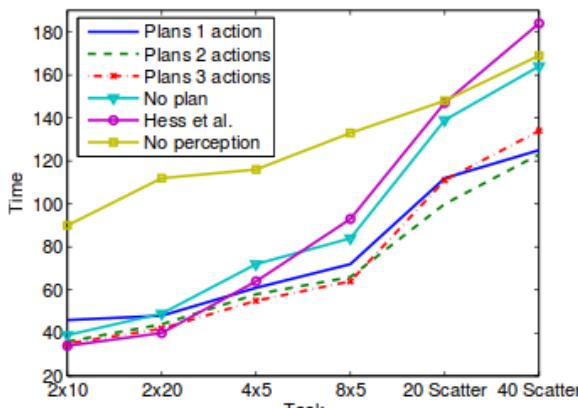
It is important highlighting that what has been done is not a benchmarking since no comparison is done with other state of the art planners, but the commented guideline for benchmarks has been used in order to get some useful metrics in order to assert the quality of the planner.

Kinds of possible experiment to do:

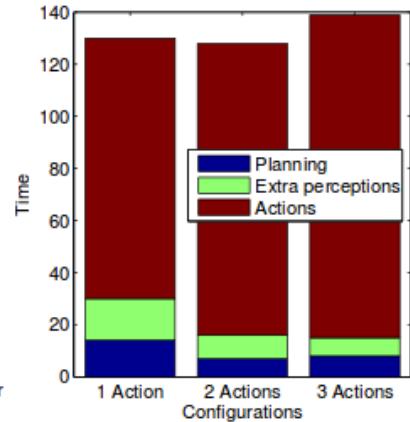
- show experiments first trying only separating the object (infinite loop behaviour) and then coupling the grasping action.
- Show the results of the planner with the pushing length chosen not accordingly to the surrounding objects but accordingly to the pushed one. From that comment the improvements.
- Show normal experiment where everything works
- Show situation why we are going to consider the risk as cost in the action.

Metrics to considers:

- number of objects
- segmentation time
- Total solution time
- planning time
- total solution
- Inverse kinematics solution time (this will probably used to justify the backtracking method)
- Predicates time computation (maybe add collision checking time in)



(a) Cleaning lentils



(b) Cleaning 40 lentils

Figure 8: (a) Skill of moving lentils to a container. The results shown are the mean over 5 runs. (b) Time for the case of moving 40 scattered lentils to a container. Extra perceptions include the movements and perceptions needed to tackle occlusions by taking the arm out of the field of view (see Section 4.2.3).

Find a way to get the computational complexity.

- 2 experiments in which everything works
- experiment which shows the limitation of the planner

7. Conclusions

7.1. Conclusions

7.2. Limitations

The planner here presented has some clear limitations, these are related to the lack of the probability and mainly to the lack of geometric information. How previously said, the pushing action is supposed to push the object far enough from the others ones, this will be true only if the pushing will be performed up to infinite, and this is not the case.. The pushing is performed taking into account the geometry of the manipulated object, without taking care about the surrounding objects geometry and poses. This problem is though to be resolved by replanning, that is, if an object has been pushed but not enough, the planner will return a sequence where the first action is again to push that object. This is actually what the planner does, but since for the most of cases, pushing along directions 1 and 2 is feasible, for the planner both directions are feasible to resolve the problem. So what it could happen is that the planner could first return as solution to push a certain object along direction 1, and then, when replanning, to push the object along direction 2. What is here commented is a infinite loop. This happens because the planner has no information between consecutive frames, that is, it consider the problem as a separated one from the previous frame. This fact is a very undesirable one and there are several scenarios that can present a behaviour like that one. The promising part is the combination of the grasping and pushing actions. It has been observed that the majority of times, although for complex scenarios, the solution is mainly based on a proper sequence of grasping action, while the pushing action is an auxiliary action which is rarely used. Taking this into account, and the presented limitation, is very likely that once an object has been pushed, ones of the objects which cannot be grasped before can now be grasped, avoiding such undesired infinite loop.

Somewhere talk about the limitation of the segmentation that is we assume it to be perfect because if not the planner will return likely an infeasible plan

7.3. Future Work

Bibliography

- [1] Fast Downward Home Page.
- [2] <http://www.openperception.org/> - open perception.
- [3] PDDL - The Planning Domain Definition Language. Technical report, CVC TR-98-003/DCS TR-1165, Yale Center for Computational Vision and Control, 1998.
- [4] Bullet <http://wiki.ros.org/bullet>, 2015.
- [5] Nichola Abdo, Henrik Kretzschmar, Luciano Spinello, and Cyrill Stachniss. Learning manipulation actions from a few demonstrations. In *Robotics and Automation (ICRA), 2013 IEEE International Conference on*, pages 1268–1275. IEEE, 2013.
- [6] Francesco Amigoni, Andrea Bonarini, Giulio Fontana, Matteo Matteucci, Viola Schiaffonati, Aamir Ahmad, Iman Awaad, Jakob Berghofer, Rainer Bischoff, et al. General evaluation criteria, modules and metrics for benchmarking through competitions. 2014.
- [7] Julien Bidot, Lars Karlsson, Fabien Lagriffoul, and Alessandro Saffiotti. Geometric backtracking for combined task and motion planning in robotic systems. *Artificial Intelligence*, 2015.
- [8] Peter Brook, Matei Ciocarlie, and Kajen Hsiao. Collaborative grasp planning with multiple object representations. In *Robotics and Automation (ICRA), 2011 IEEE International Conference on*, pages 2851–2858. IEEE, 2011.
- [9] Rui Coelho and Alexandre Bernardino. Planning push and grasp actions: Experiments on the icub robot.
- [10] Nikolaus Correll, Kostas E. Bekris, Dmitry Berenson, Oliver Brock, Albert Causo, Kris Hauser, Kei Okada, Alberto Rodriguez, Joseph M. Romano, and Peter R. Wurman. Lessons from the amazon picking challenge. *CoRR*, abs/1601.05484, 2016.
- [11] N. Covallero and G. Alenyà. Grasping novel objects. Technical Report IRI-TR-16-01, Institut de Robòtica i Informàtica Industrial, CSIC-UPC, 2016.
- [12] Richard Dearden and Chris Burbridge. An approach for efficient planning of robotic manipulation tasks. In *ICAPS*. Citeseer, 2013.
- [13] Richard Dearden and Chris Burbridge. Manipulation planning using learned symbolic state abstractions. *Robotics and Autonomous Systems*, 62(3):355 – 365, 2014. Advances in Autonomous Robotics — Selected extended papers of the joint 2012 {TAROS} Conference and the {FIRA} RoboWorld Congress, Bristol, {UK}.
- [14] Mehmet Dogar and Siddhartha Srinivasa. A framework for push-grasping in clutter. In Nick Roy Hugh Durrant-Whyte and Pieter Abbeel, editors, *Robotics: Science and Systems VII*. MIT Press, July 2011.

- [15] Mustafa Ersen, Melodi Deniz Ozturk, Mehmet Biberci, Sanem Sariel, and Hulya Yalcin. Scene interpretation for lifelong robot learning. In *Proceedings of the 9th international workshop on cognitive robotics (CogRob 2014) held in conjunction with ECAI-2014*, 2014.
- [16] Mustafa Ersen, Sanem Sariel Talay, and Hulya Yalcin. Extracting spatial relations among objects for failure detection. In *KIK@ KI*, pages 13–20, 2013.
- [17] David Fischinger, Astrid Weiss, and Markus Vincze. Learning grasps with topographic features. volume 34, pages 1167–1194, 2015.
- [18] Malte Helmert. The fast downward planning system. *J. Artif. Intell. Res.(JAIR)*, 26:191–246, 2006.
- [19] Tucker Hermans, James M. Rehg, and Aaron F. Bobick. Guided pushing for object singulation. In *IROS*, pages 4783–4790. IEEE, 2012.
- [20] Ian Jolliffe. *Principal component analysis*. Wiley Online Library, 2002.
- [21] Dov Katz. *Interactive perception of articulated objects for autonomous manipulation*. University of Massachusetts Amherst, 2011.
- [22] Dov Katz, Moslem Kazemi, J. Andrew (Drew) Bagnell, and Anthony (Tony) Stentz . Clearing a pile of unknown objects using interactive perception. In *Proceedings of IEEE International Conference on Robotics and Automation*, March 2013.
- [23] Dov Katz, Arun Venkatraman, Moslem Kazemi, J Andrew Bagnell, and Anthony Stentz. Perceiving, learning, and exploiting object affordances for autonomous pile manipulation. *Autonomous Robots*, 37(4):369–382, 2014.
- [24] Thomas Keller and Patrick Eyerich. Prost: Probabilistic planning based on uct. In *ICAPS*, 2012.
- [25] Andrey Kolobov, Mausam, and Daniel S Weld. Lrtdp vs. uct for online probabilistic planning. In *Twenty-Sixth AAAI Conference on Artificial Intelligence*, 2012.
- [26] Steven M. LaValle. *Planning Algorithms*. Cambridge University Press, May 2006.
- [27] Iain Little, Sylvie Thiebaux, et al. Probabilistic planning vs. replanning. In *ICAPS Workshop on IPC: Past, Present and Future*, 2007.
- [28] Bhaskara Marthi, Stuart J Russell, and Jason Wolfe. Angelic semantics for high-level actions. In *ICAPS*, pages 232–239, 2007.
- [29] David Martínez, Guillem Alenyà, and Carme Torras. Planning robot manipulation to clean planar surfaces. *Engineering Applications of Artificial Intelligence*, 39:23–32, 2015.
- [30] Lorenz Mösenlechner and Michael Beetz. Using physics- and sensor-based simulation for high-fidelity temporal projection of realistic robot behavior. In *AIPS*, 2009.
- [31] Dana Nau, Yue Cao, Amnon Lotem, and Hector Munoz-Avila. Shop: Simple hierarchical ordered planner. In *Proceedings of the 16th international joint conference on Artificial intelligence-Volume 2*, pages 968–973. Morgan Kaufmann Publishers Inc., 1999.

- [32] Melodi Ozturk, Mustafa Ersen, Melis Kapotoglu, Cagatay Koc, Sanem Sariel-Talay, and Hulya Yalcin. Scene interpretation for self-aware cognitive robots. 2014.
- [33] Jia Pan, Sachin Chitta, and Dinesh Manocha. Fcl: A general purpose library for collision and proximity queries. In *Robotics and Automation (ICRA), 2012 IEEE International Conference on*, pages 3859–3866. IEEE, 2012.
- [34] Jeremie Papon, Alexey Abramov, Markus Schoeler, and Florentin Wörgötter. Voxel cloud connectivity segmentation - supervoxels for point clouds. In *Computer Vision and Pattern Recognition (CVPR), 2013 IEEE Conference on*, Portland, Oregon, June 22-27 2013.
- [35] Stuart J. Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Pearson Education, 2 edition, 2003.
- [36] S. C. Stein, M. Schoeler, J. Papon, and F. Woergoetter. Object partitioning using local convexity. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2014.
- [37] Andreas ten Pas and Robert Platt. Using geometry to detect grasp poses in 3d point clouds. In *International Symposium on Robotics Research (ISRR)*, September 2015.
- [38] Nikolaus Vahrenkamp, Tamim Asfour, and Rudiger Dillmann. Robot placement based on reachability inversion. In *Robotics and Automation (ICRA), 2013 IEEE International Conference on*, pages 1970–1975. IEEE, 2013.